

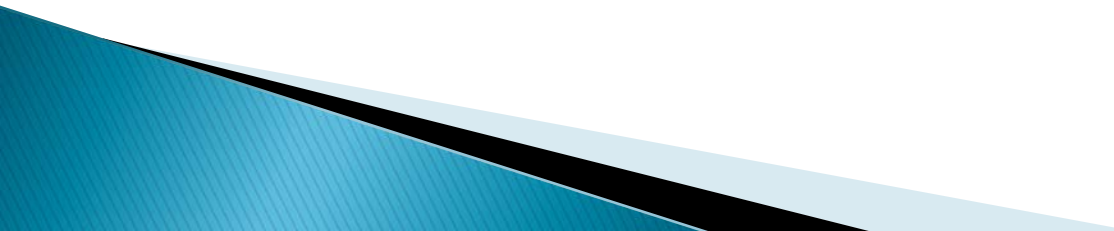
# CMSC 726

## Lecture 13: Evaluation

Lise Getoor  
October 14, 2010

**ACKNOWLEDGEMENTS:** The material in this course is a synthesis of materials from many sources, including: Hal Daume III, Tom Dietterich, Mark Drezde, Carlos Guestrin, **Andrew Moore**, Andrew Ng, Ben Taskar, Eric Xing, and others. I am very grateful for their generous sharing of insights and materials.

# Outline

- ▶ Evaluating Hypotheses & Learning Algorithms
    - Performance Criteria
    - Motivation
    - Cross Validation
    - Confidence Intervals
  - ▶ Bias–Variance Decomposition for Regression
  - ▶ Summary and Conclusion
- 

# Evaluating Models

- ▶ Need a measure of value – the cost (loss, utility) of a model
- ▶ Often use accuracy (or error)
  - Accuracy – how many examples we get “right”
  - Error – how many examples we get wrong
- ▶ Can be weighted
  - If examples are not equal, could count the cost (or utility) of mispredicted (correct) examples

# Supervised Learning Performance Criteria

- ▶ Accuracy/Squared Error/Probability Calibration
- ▶ Others:
  - Area under the ROC Curve
  - Lift
  - F-Score
  - Average Precision
  - Precision/Recall Break-Even Point
  - Confusion Matrixes

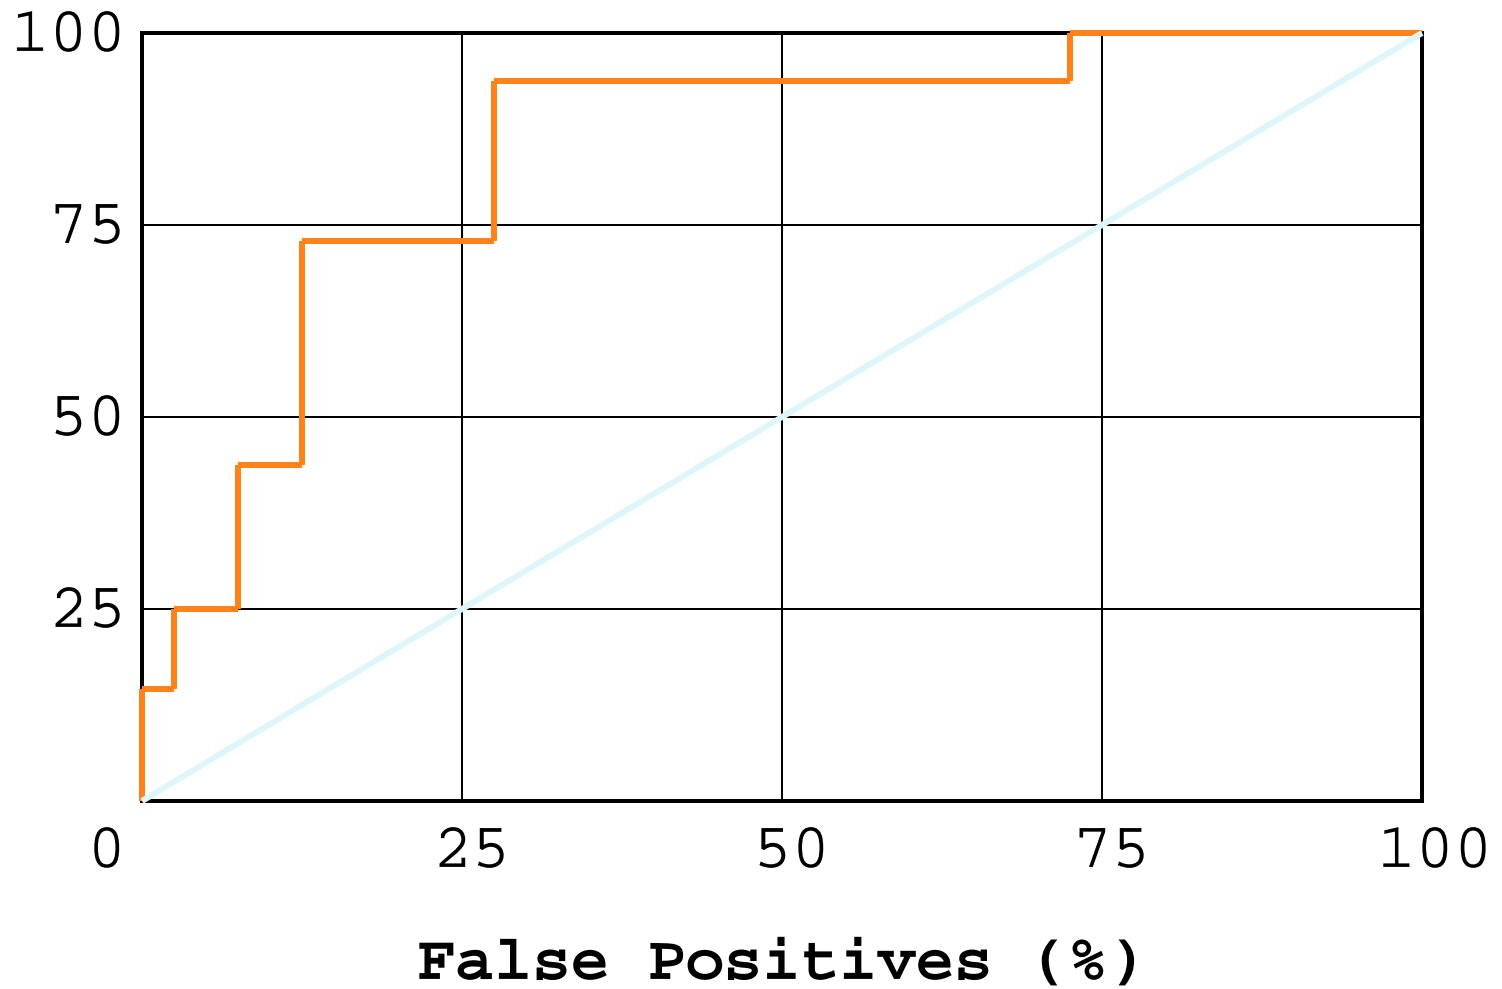
# Types of Performance Criteria for Supervised Learning

- ▶ Interpret as threshold
  - classification accuracy
- ▶ Interpret as probabilities
  - (conditional) likelihood, squared error
- ▶ Interpret as ranking
  - ROC curves

# Receiver Operator Characteristic (ROC) Curves

- ▶ Originally from signal detection
- ▶ Becoming very popular for ML
- ▶ Used in:
  - Two class problems
  - Where predictions are ordered in some way (e.g., neural network activation is often taken as an indication of how strong or weak a prediction is)
- ▶ Plotting an ROC curve:
  - Sort predictions by their predicted strength
  - Start at the bottom left
  - For each positive example, go up  $1/P$  units where  $P$  is the number of positive examples
  - For each negative example, go right  $1/N$  units where  $N$  is the number of negative examples

# ROC Curve



# ROC Properties

- ▶ Can visualize the tradeoff between coverage and accuracy (as we lower the threshold for prediction how many more true positives will we get in exchange for more false positives)
- ▶ Gives a better feel when comparing algorithms
  - Algorithms may do well in different portions of the curve
- ▶ A perfect curve would start in the bottom left, go to the top left, then over to the top right
  - A random prediction curve would be a line from the bottom left to the top right
- ▶ When comparing curves:
  - Can look to see if one curve dominates the other (is always better)
  - Can compare the area under the curve (very popular – some people even do t-tests on these numbers)



# Lift

- ▶ Lift measures how much better a classifier is at predicting positives than a baseline classifier that randomly predicts positives (at the rate observed for positives in the data)

$$\text{LIFT} = \frac{\% \text{ of true positives about the threshold}}{\% \text{ of dataset about the threshold}}$$

# Precision/Recall

- ▶ Precision: fraction of examples predicted as positive that are actually positive
- ▶ Recall: fraction of true positives that are predicted as positives
- ▶ Combining measures:
  - precision–recall F score: harmonic mean of the precision and recall at a given threshold
  - precision at recall level: set recall, measure precision
  - break even point: the precision at which the precision equals recall
  - average precision: average of the precisions at eleven evenly spaced recall levels.

# Confusion Matrix

		Predicted		
		Positive	Negative	Total
Actual	Positive	True Positive (TP)	False Negative (FN)	#Positives
	Negative	False Positive (FP)	True Negative (TN)	#Negatives
	Total	TP+FP	FN+TN	#Examples

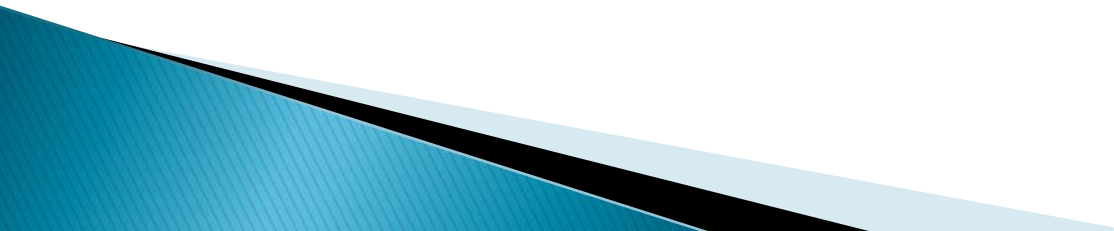
- ▶  $\text{Accuracy} = (\text{TP} + \text{TN}) / \text{\#Examples}$
- ▶  $\text{Error} = (\text{FP} + \text{FN}) / \text{\#Examples}$
- ▶  $\text{Recall (sensitivity, true positive rate)} = \text{TP} / \text{\#Positives}$
- ▶  $\text{Precision} = \text{TP} / (\text{FP} + \text{TP})$
- ▶  $\text{True Negative Rate (specificity)} = \text{TN} / \text{\#Negatives}$
- ▶  $\text{False Positive Rate} = \text{FP} / (\text{FP} + \text{TP})$
- ▶  $\text{False Negative Rate} = \text{FN} / \text{\#Negatives}$

# Confusion Matrix – Multi Class

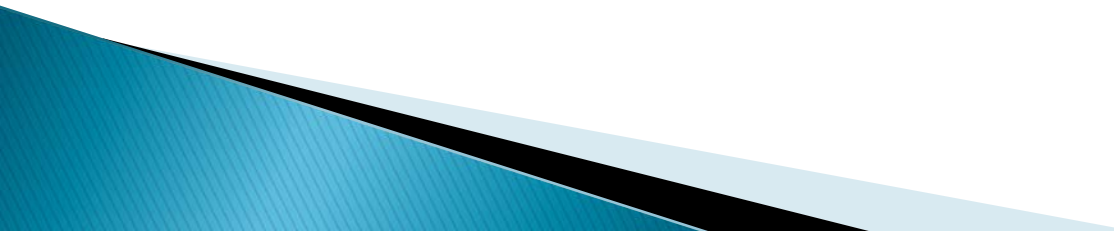
- ▶ For many problems (especially multiclass problems), often useful to examine the sources of error
- ▶ Confusion matrix:

		Predicted			Total
		ClassA	ClassB	ClassC	
Expected	ClassA	25	5	20	50
	ClassB	0	45	5	50
	ClassC	25	0	25	50
	Total	50	50	50	150

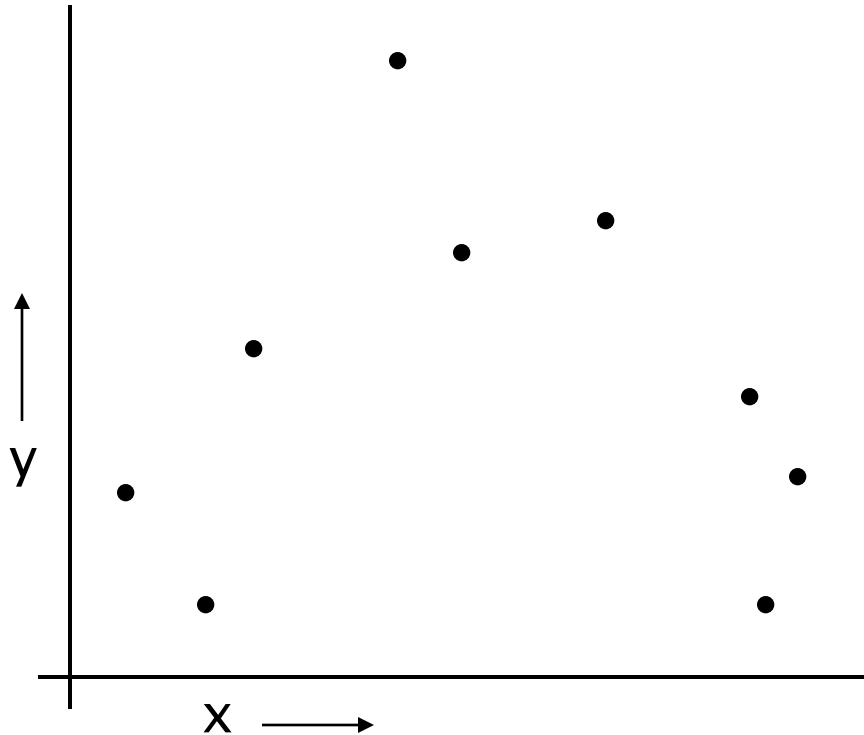
# Results Analysis: Confusion Matrix

- ▶ Building a confusion matrix
    - Zero all entries
    - For each data point add one in row corresponding to actual class of problem under column corresponding to predicted class
  - ▶ Perfect prediction has all non-zeros values down the diagonal
  - ▶ Off diagonal entries can often tell us about what is being mispredicted
- 

# Outline

- ▶ Evaluating Hypotheses & Learning Algorithms
    - Performance Criteria
    - **Motivation**
    - Cross Validation
    - Confidence Intervals
  - ▶ Bias–Variance Decomposition for Regression
  - ▶ Summary and Conclusion
- 

# Motivation: A Regression Problem

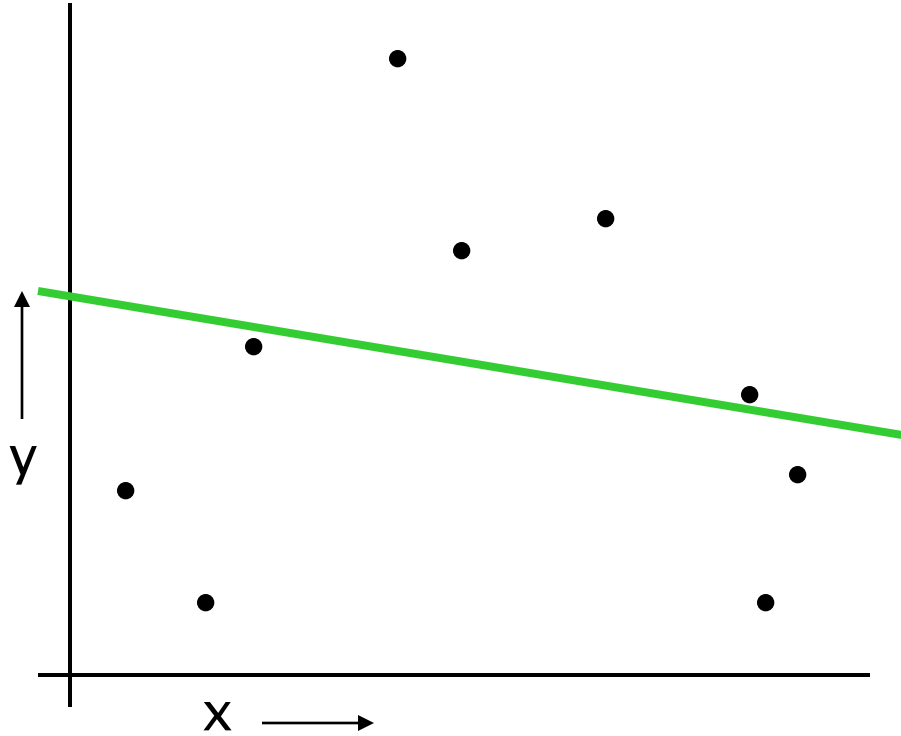


$$y = f(x) + \text{noise}$$

Can we learn  $f$  from this data?

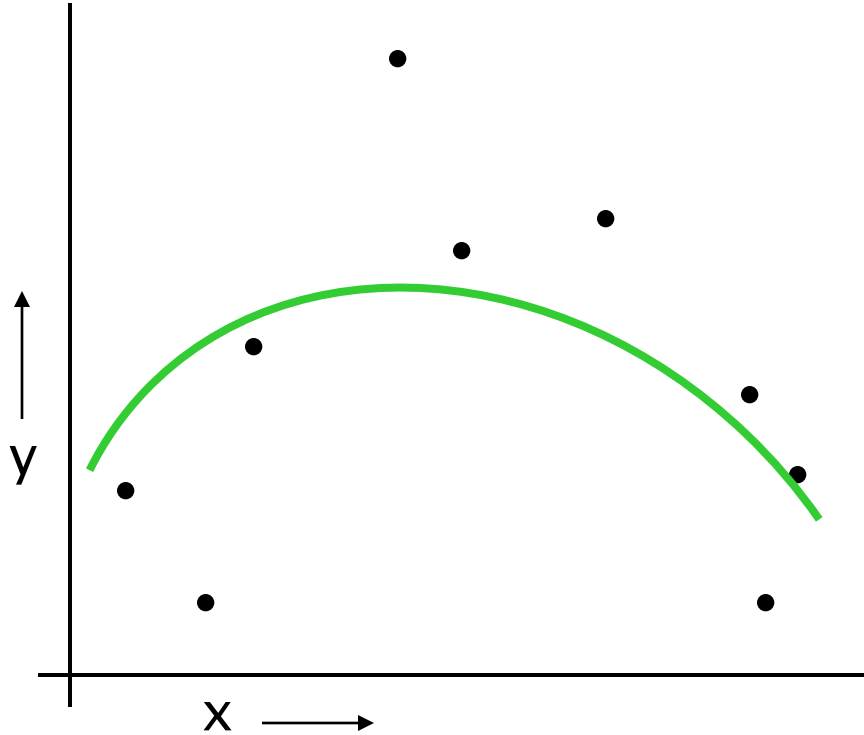
Let's consider three methods...

# Linear Regression

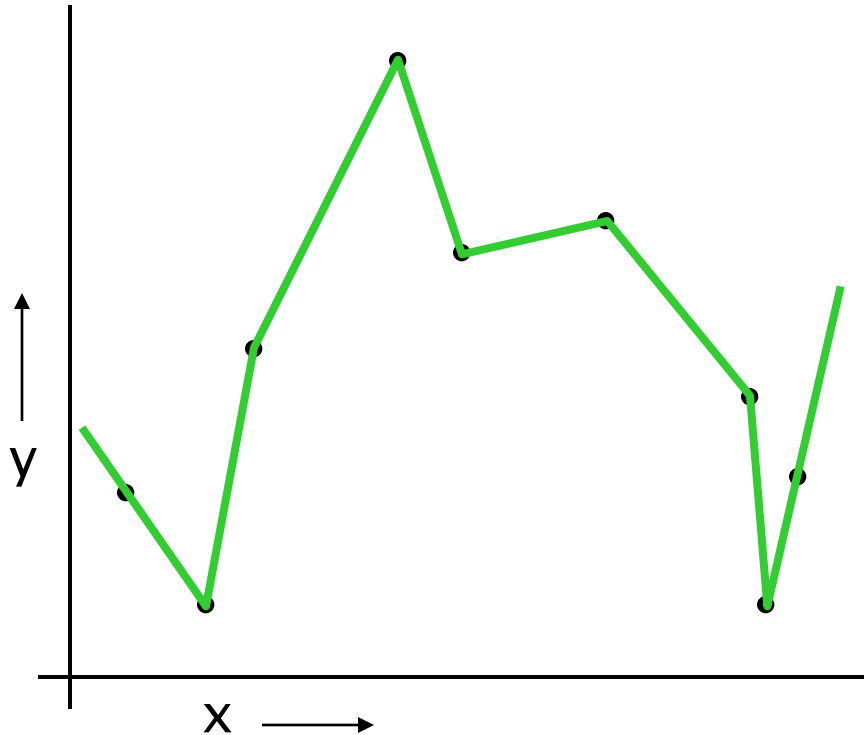




# Quadratic Regression

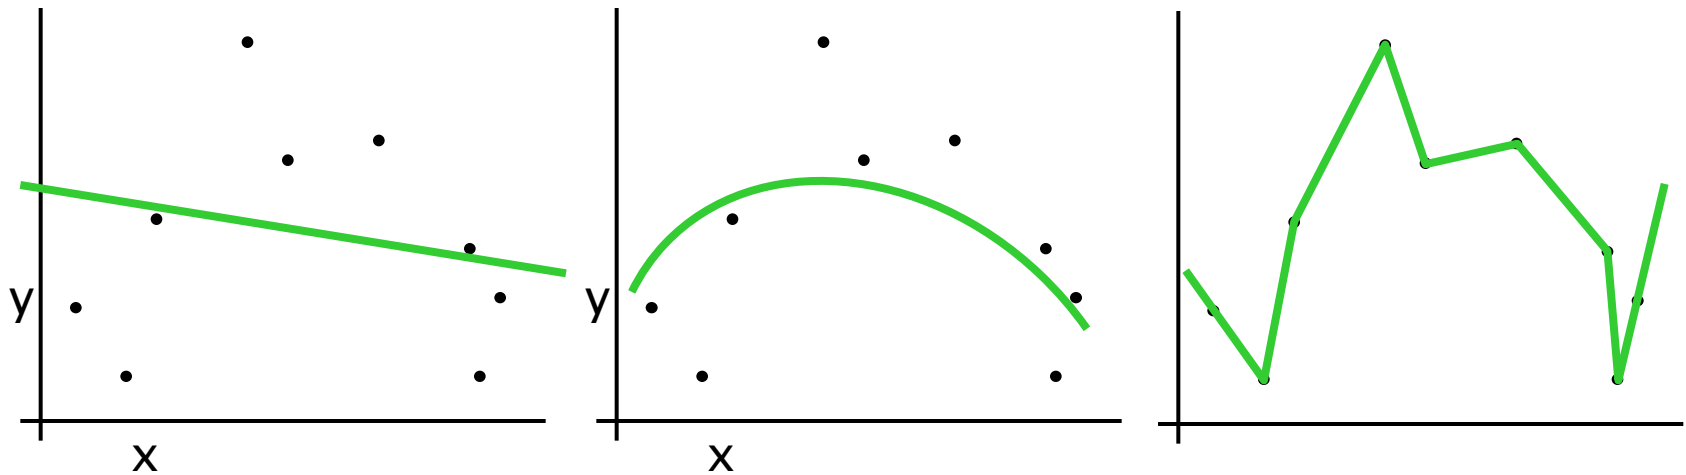


# Join-the-dots



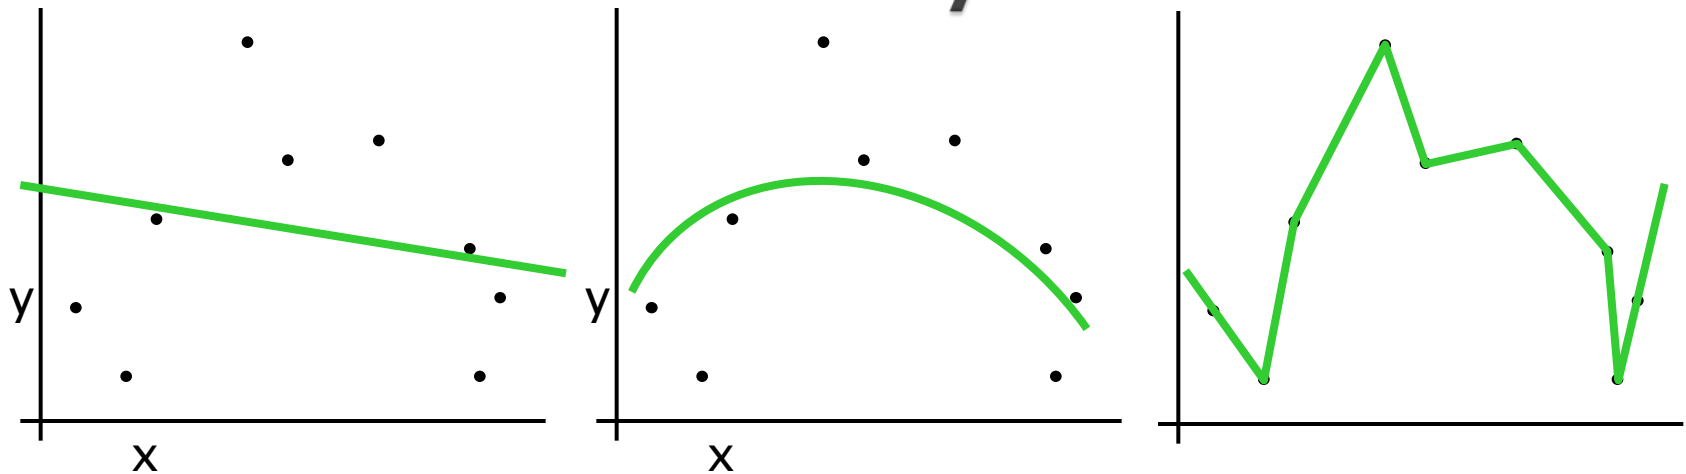
Also known as **piecewise linear nonparametric regression** if you think that sounds better

# Which is best?



Why not choose the method with the best fit to the data?

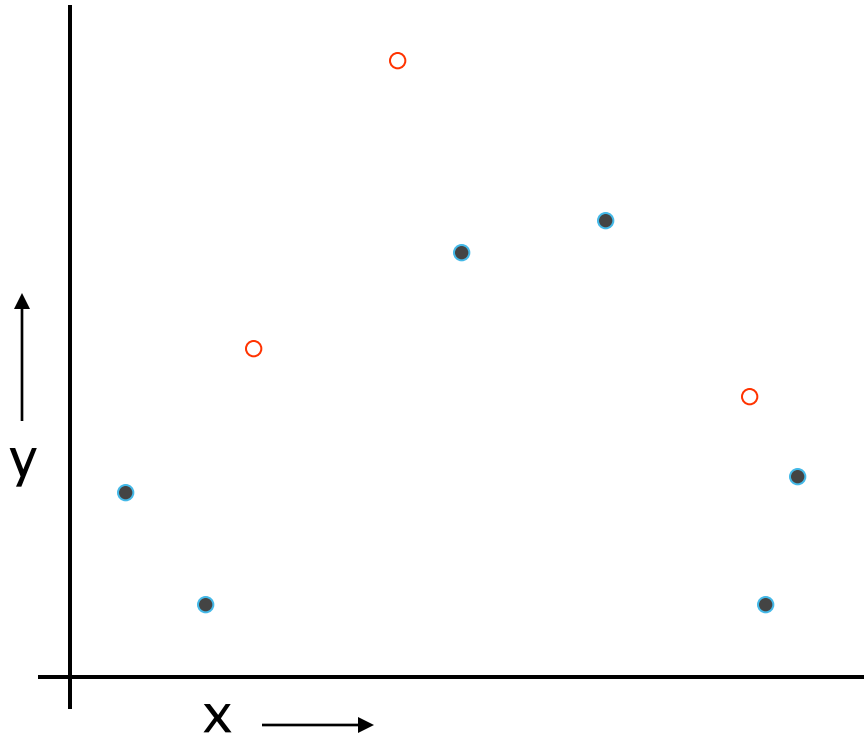
# What do we really want?



Why not choose the method with the best fit to the data?

“How well are you going to **predict** future data drawn from the same distribution?”

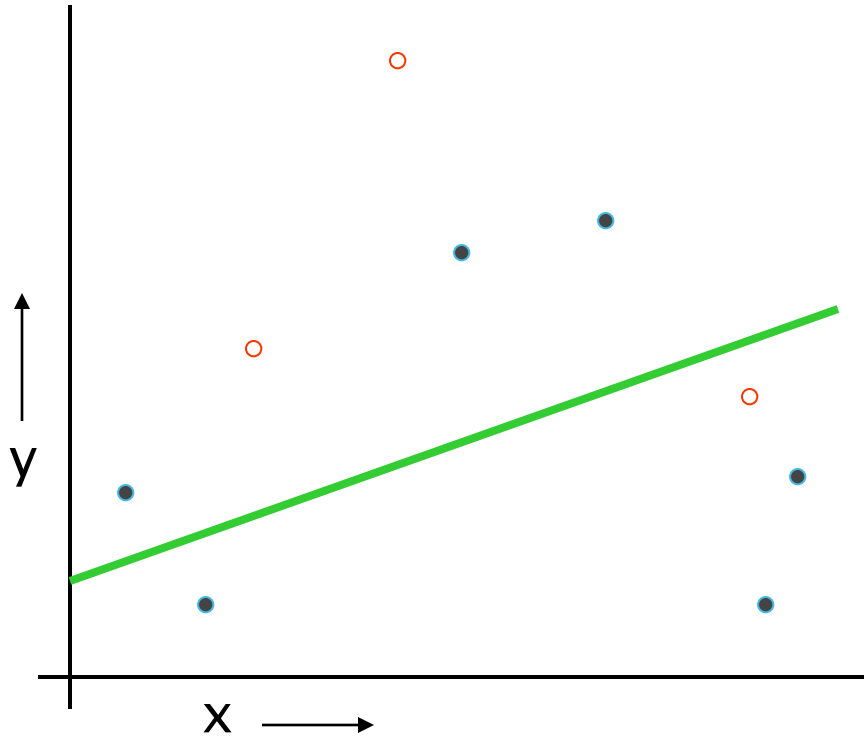
# The test set method



1. Randomly choose 30% of the data to be in a **test set**

2. The remainder is a training set

# The test set method



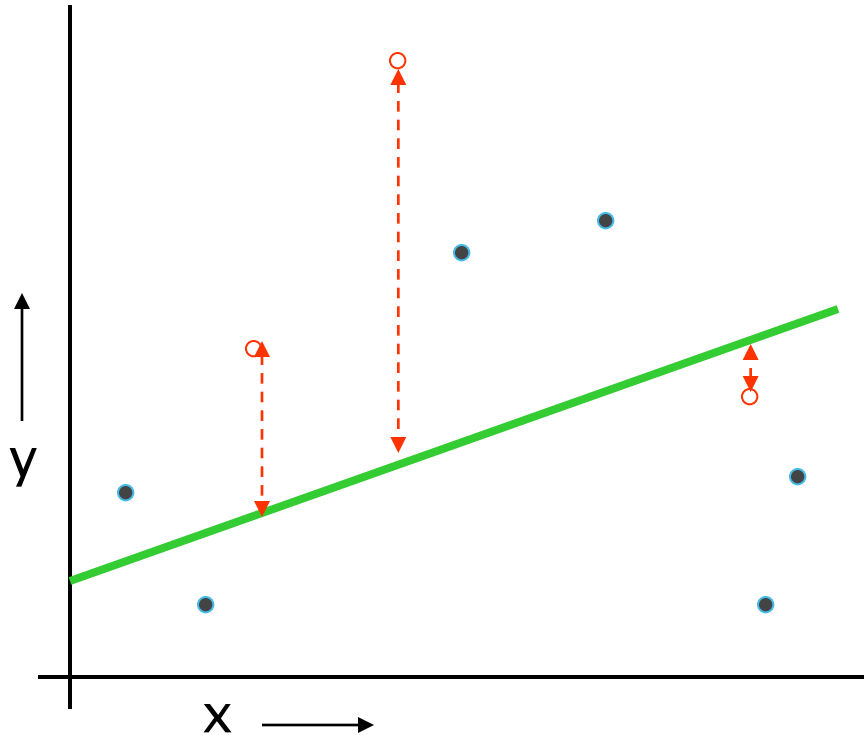
1. Randomly choose 30% of the data to be in a **test set**

2. The remainder is a training set

3. Perform your regression on the training set

(Linear regression example)

# The test set method



(Linear regression example)

Mean Squared Error = 2.4

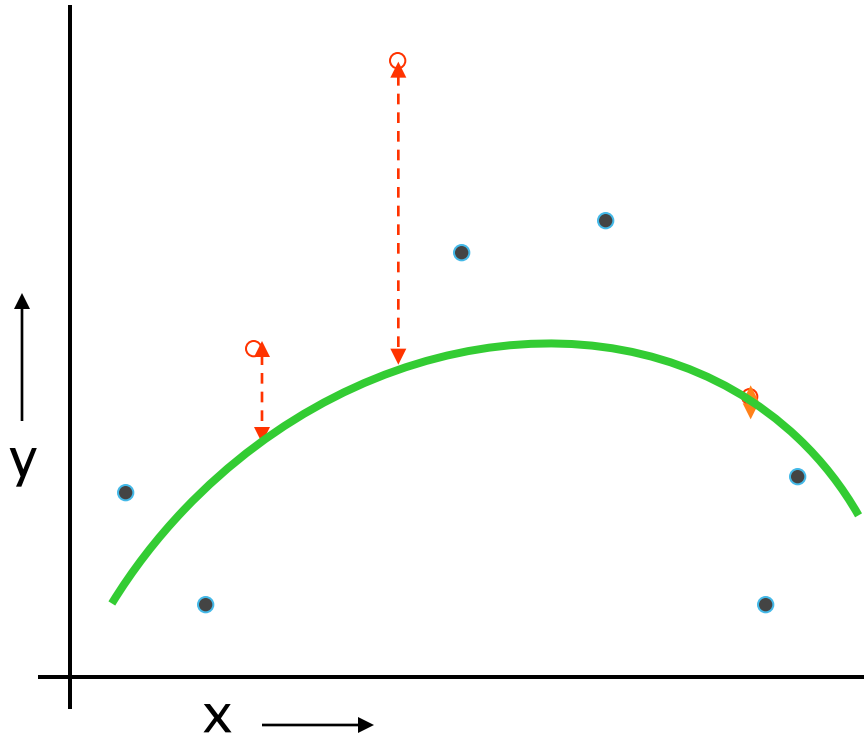
1. Randomly choose 30% of the data to be in a **test set**

2. The remainder is a training set

3. Perform your regression on the training set

4. Estimate your future performance with the **test set**

# The test set method



(Quadratic regression example)

Mean Squared Error = 0.9

1. Randomly choose 30% of the data to be in a **test set**

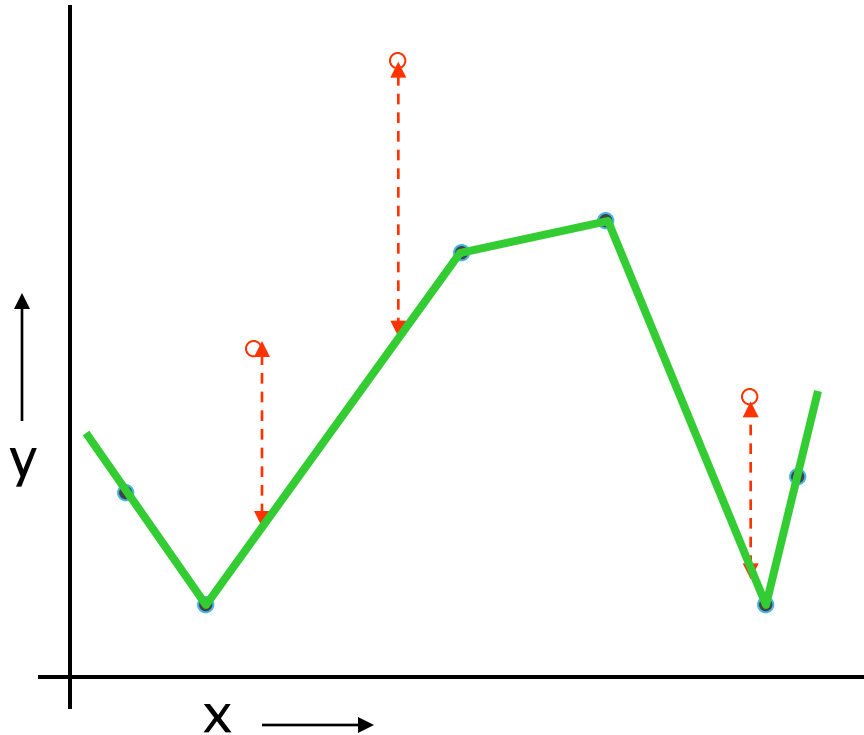
2. The remainder is a training set

3. Perform your regression on the training set

4. Estimate your future performance with the **test set**



# The test set method



(Join the dots example)

Mean Squared Error = 2.2

1. Randomly choose 30% of the data to be in a **test set**

2. The remainder is a training set

3. Perform your regression on the training set

4. Estimate your future performance with the **test set**

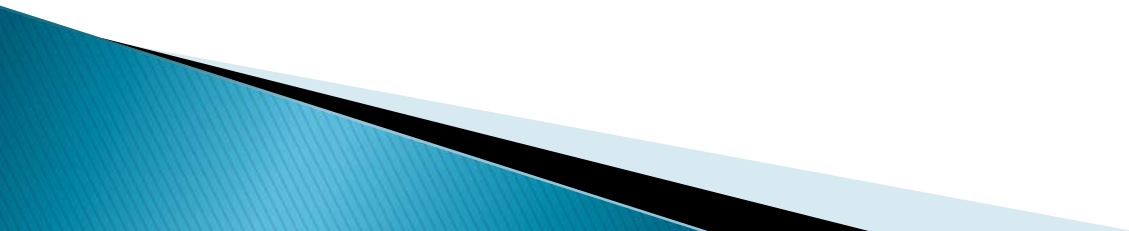
# The test set method

## Good news:

- Very very simple
- Can then simply choose the method with the best test-set score

## Bad news:

- What's the downside?



# The test set method

## Good news:

- Very very simple
- Can then simply choose the method with the best test-set score

## Bad news:

- Wastes data: we get an estimate of the best method to apply to 30% less data
- If we don't have much data, our test-set might just be lucky or unlucky

We say the “test-set estimator of performance has high variance”

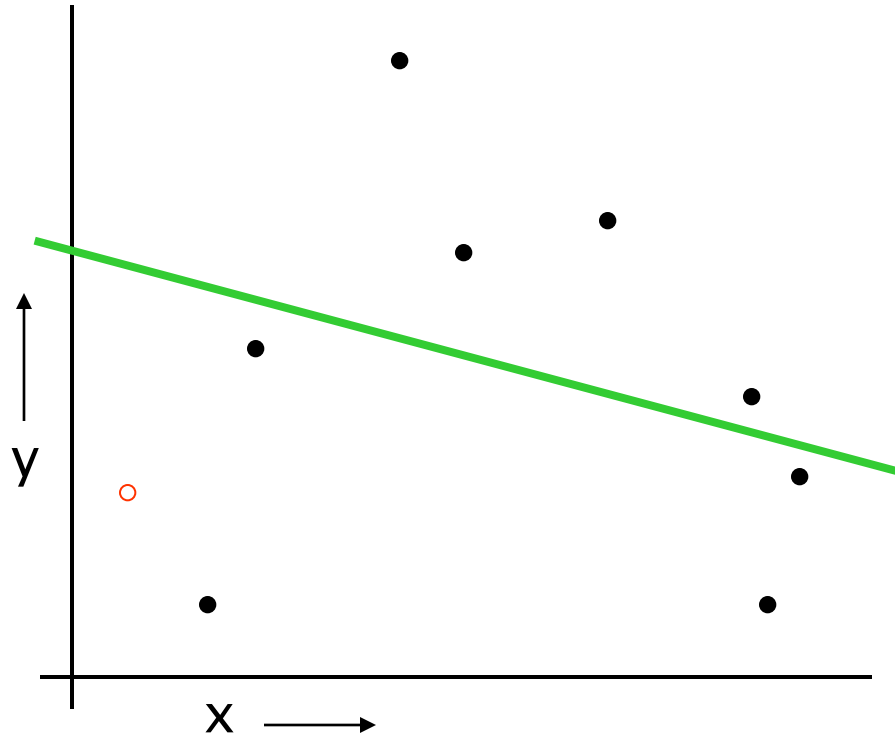
# N-Fold Cross Validation

- ▶ Popular testing methodology
- ▶ Divide data into  $N$  even-sized random folds
- ▶ For  $n = 1$  to  $N$ 
  - Train set = all folds except  $n$
  - Test set = fold  $n$
  - Create learner with train set
  - Count number of errors on test set
- ▶ Accumulate number of errors across  $N$  test sets and divide by  $N$  (result is error rate)
- ▶ For comparing algorithms, use the same set of folds to create learners (results are paired)

# N-Fold Cross Validation

- ▶ Advantages/disadvantages
  - Estimate of error within a single data set
  - Every point used once as a test point
  - At the extreme (when  $N = \text{size of data set}$ ), called leave-one-out testing

# LOOCV (Leave-one-out Cross Validation)

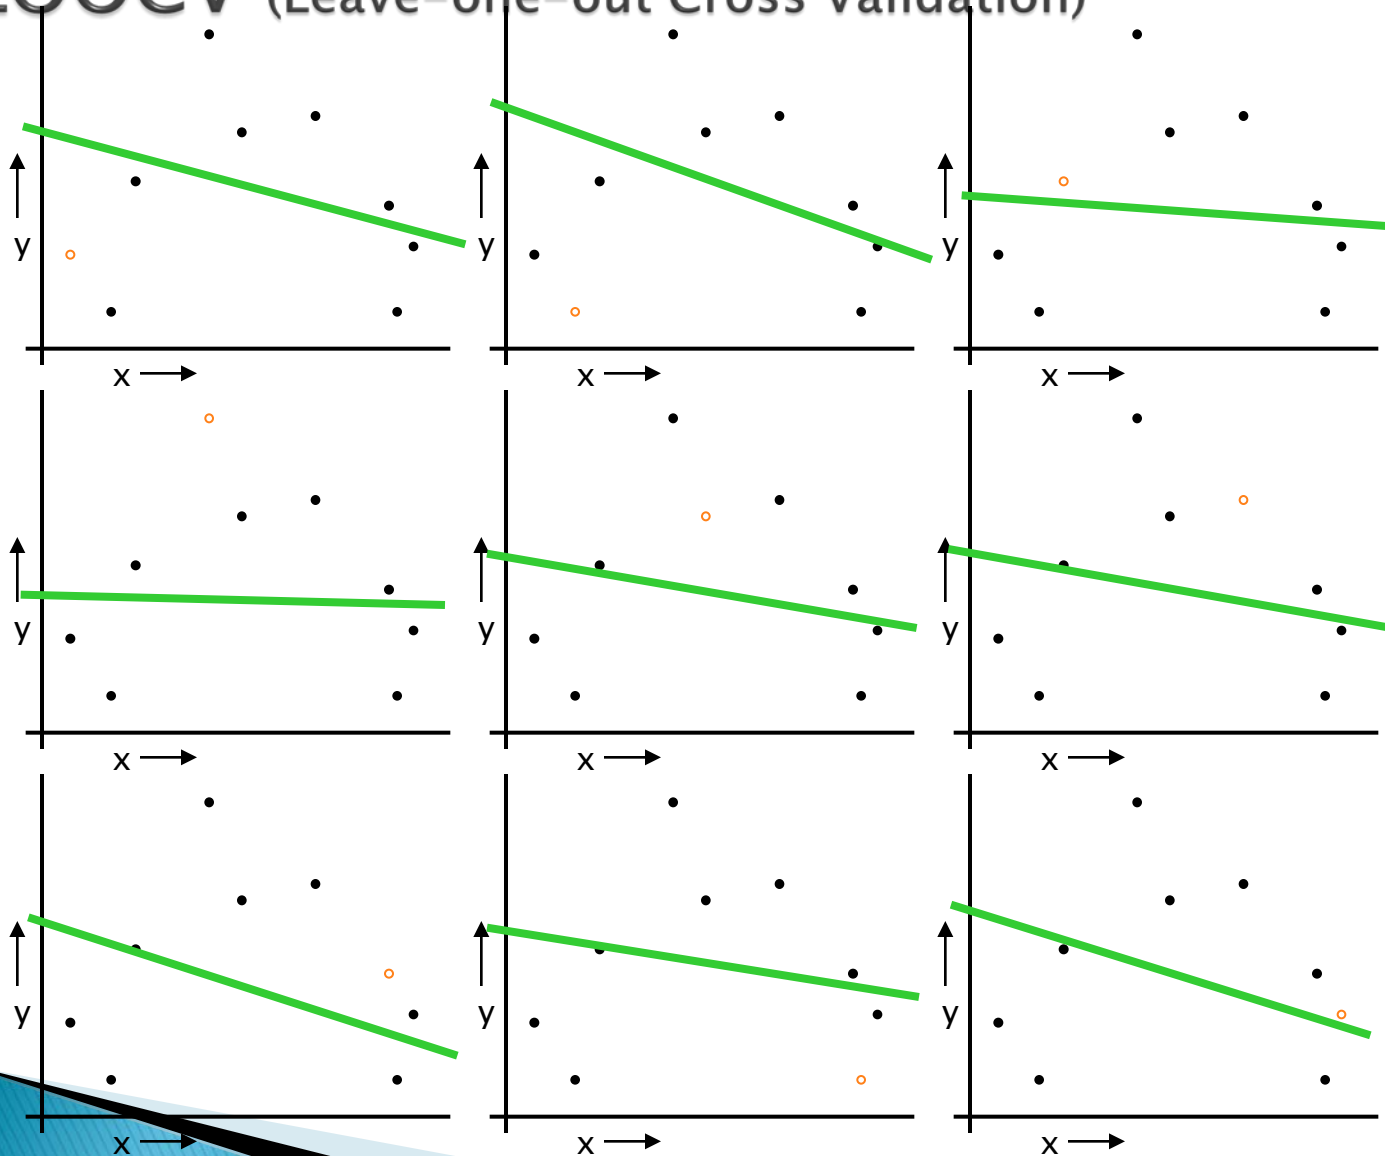


For  $k=1$  to  $N$

1. Let  $(x_k, y_k)$  be the  $k^{\text{th}}$  record
2. Temporarily remove  $(x_k, y_k)$  from the dataset
3. Train on the remaining  $N-1$  datapoints
4. Note your error  $(x_k, y_k)$

When you've done all points, report the mean error.

# LOOCV (Leave-one-out Cross Validation)



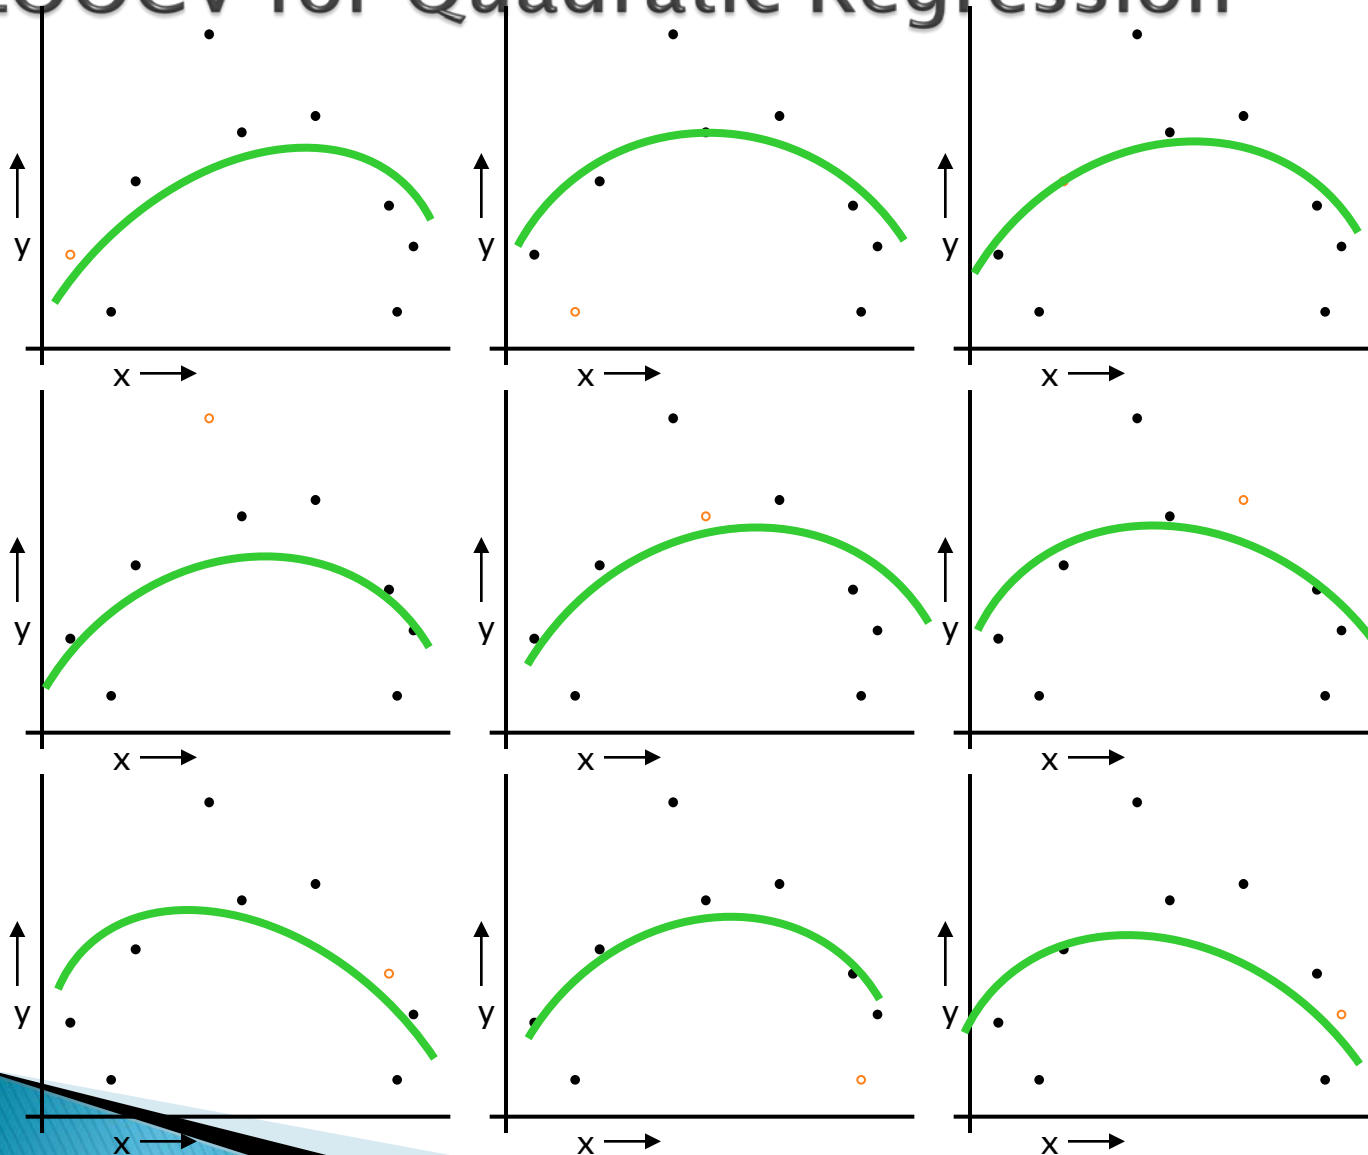
For  $k=1$  to  $R$

1. Let  $(x_k, y_k)$  be the  $k^{\text{th}}$  record
2. Temporarily remove  $(x_k, y_k)$  from the dataset
3. Train on the remaining  $R-1$  datapoints
4. Note your error  $(x_k, y_k)$

When you've done all points, report the mean error.

$$MSE_{LOOCV} = 2.12$$

# LOOCV for Quadratic Regression



For  $k=1$  to  $R$

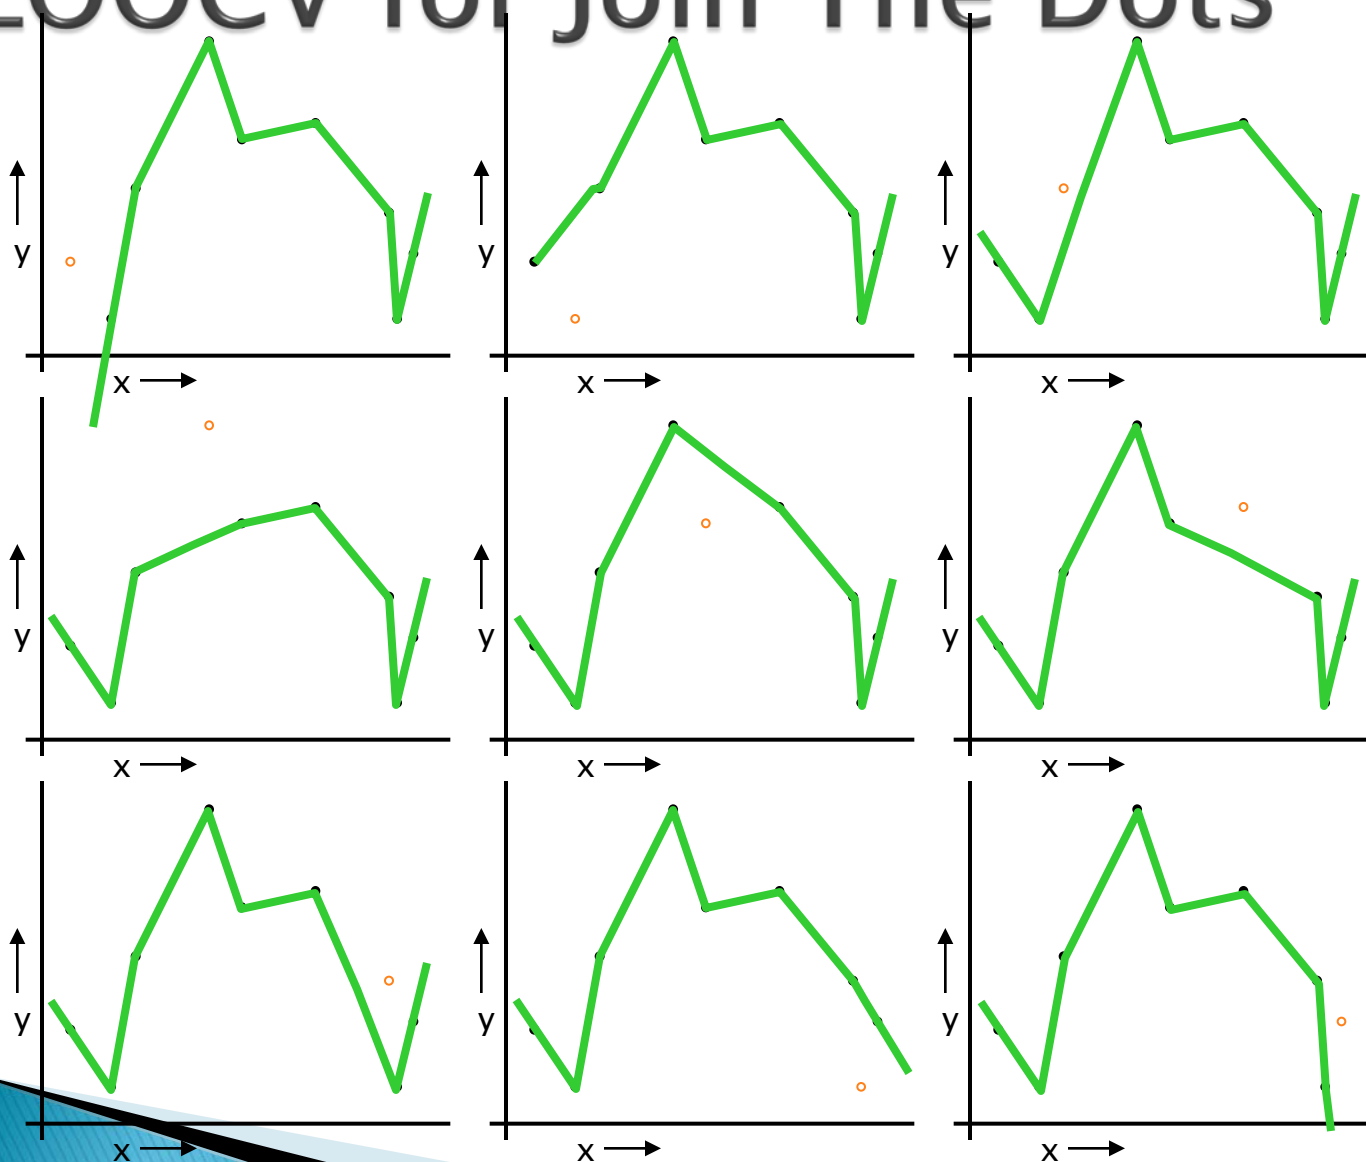
1. Let  $(x_k, y_k)$  be the  $k^{\text{th}}$  record
2. Temporarily remove  $(x_k, y_k)$  from the dataset
3. Train on the remaining  $R-1$  datapoints
4. Note your error  $(x_k, y_k)$

When you've done all points, report the mean error.

$$MSE_{LOOCV} = 0.962$$



# LOOCV for Join The Dots



For  $k=1$  to  $R$

1. Let  $(x_k, y_k)$  be the  $k^{\text{th}}$  record
2. Temporarily remove  $(x_k, y_k)$  from the dataset
3. Train on the remaining  $R-1$  datapoints
4. Note your error  $(x_k, y_k)$

When you've done all points, report the mean error.

$$MSE_{LOOCV} = 3.33$$

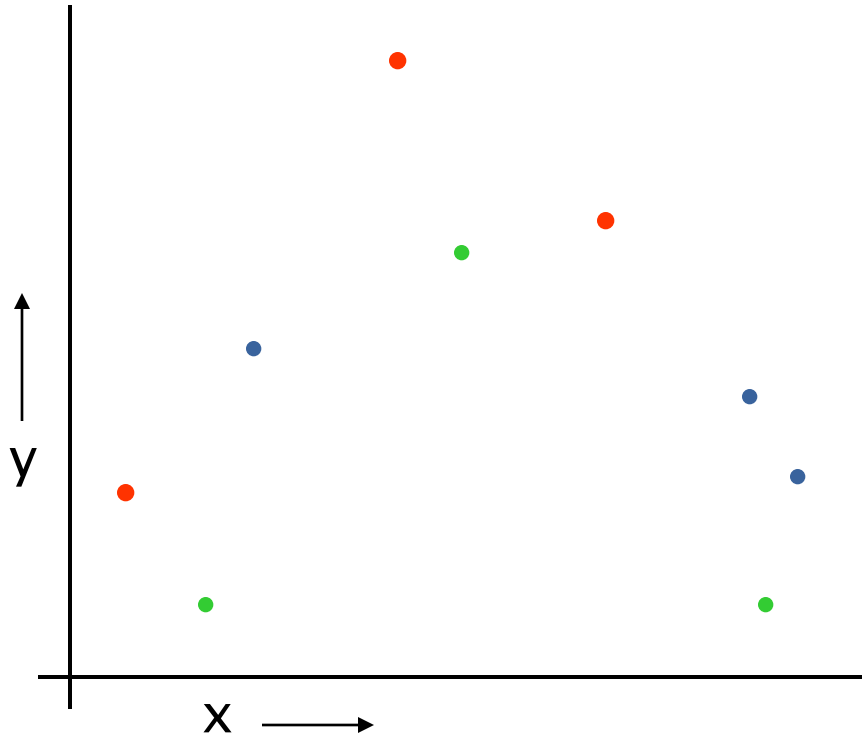
# Which kind of Cross Validation?

	<b>Downside</b>	<b>Upside</b>
<b>Test-set</b>	Variance: unreliable estimate of future performance	Cheap
<b>Leave-one-out</b>	Expensive.	Doesn't waste data

..can we get the best of both worlds?

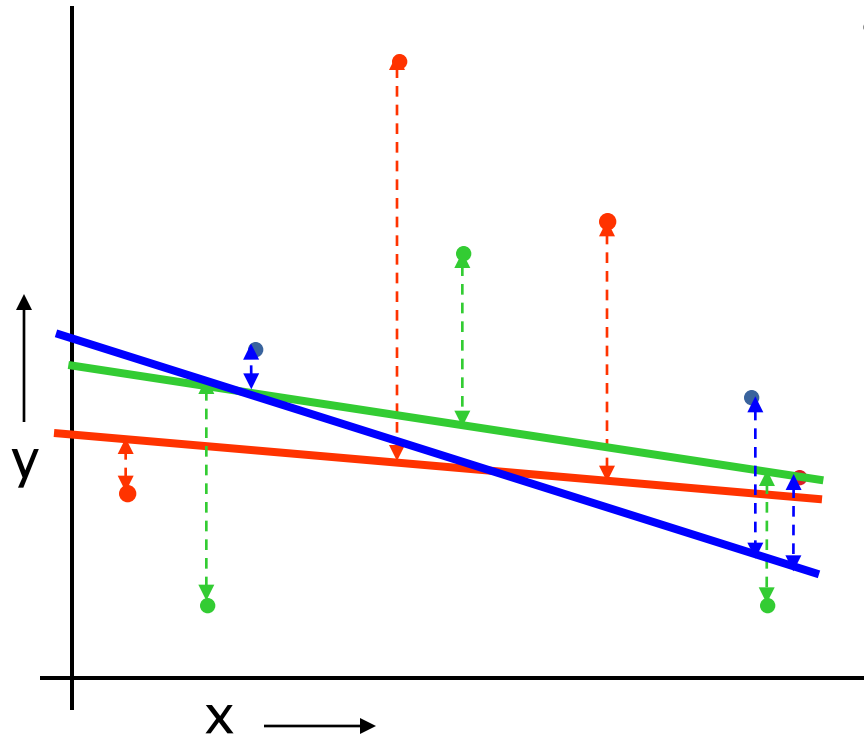


# k-fold CV



Randomly break the dataset into  $k$  partitions (in our example we'll have  $k=3$  partitions colored Red Green and Blue)

# k-fold CV



Linear Regression  
 $MSE_{3FOLD} = 2.05$

Randomly break the dataset into k partitions (in our example we'll have  $k=3$  partitions colored Red Green and Blue)

For the red partition: Train on all the points not in the red partition. Find the test-set sum of errors on the red points.

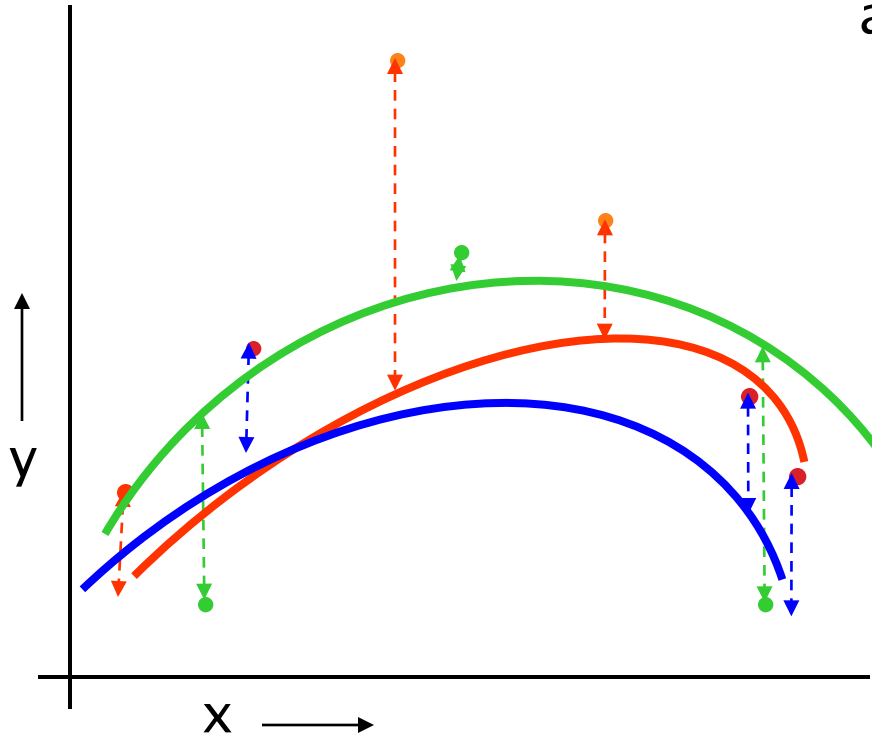
For the green partition: Train on all the points not in the green partition. Find the test-set sum of errors on the green points.

For the blue partition: Train on all the points not in the blue partition. Find the test-set sum of errors on the blue points.

Then report the mean error

# k-fold CV

Randomly break the dataset into k partitions (in our example we'll have  $k=3$  partitions colored Red Green and Blue)



Quadratic Regression  
 $MSE_{3FOLD} = 1.11$

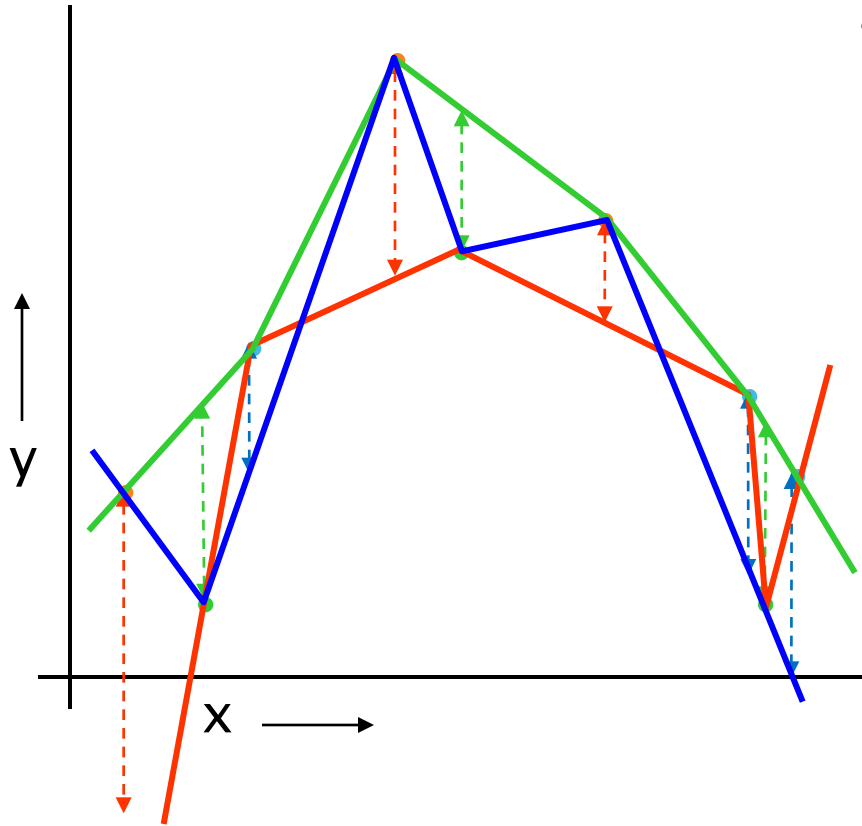
For the red partition: Train on all the points not in the red partition. Find the test-set sum of errors on the red points.

For the green partition: Train on all the points not in the green partition. Find the test-set sum of errors on the green points.

For the blue partition: Train on all the points not in the blue partition. Find the test-set sum of errors on the blue points.

Then report the mean error

# k-fold CV



Randomly break the dataset into  $k$  partitions (in our example we'll have  $k=3$  partitions colored Red Green and Blue)

For the red partition: Train on all the points not in the red partition. Find the test-set sum of errors on the red points.

For the green partition: Train on all the points not in the green partition. Find the test-set sum of errors on the green points.

For the blue partition: Train on all the points not in the blue partition. Find the test-set sum of errors on the blue points.

Joint-the-dots  
 $MSE_{3FOLD} = 2.93$

Then report the mean error

# Which kind of Cross Validation?

	<b>Downside</b>	<b>Upside</b>
<b>Test-set</b>	<b>Variance: unreliable estimate of future performance</b>	<b>Cheap</b>
<b>Leave-one-out</b>	<b>Expensive.</b>	<b>Doesn't waste data</b>
<b>10-fold</b>	<b>Wastes 10% of the data. 10 times more expensive than test set</b>	<b>Only wastes 10%. Only 10 times more expensive instead of N times.</b>
<b>3-fold</b>	<b>Wastes more than 10-fold. more expensive than test set</b>	<b>Slightly better than test-set</b>
<b>N-fold</b>	<b>Identical to Leave-one-out</b>	

# Two Definitions of Error

The **true error** of hypothesis  $h$  with respect to target function  $f$  and distribution  $D$  is the probability that  $h$  will misclassify an instance drawn at random according to  $D$ .

$$error_D(h) \equiv \Pr_{x \in D}[f(x) \neq h(x)]$$

The **sample error** of  $h$  with respect to target function  $f$  and data sample  $S$  is the proportion of examples  $h$  misclassifies

$$error_S(h) \equiv \frac{1}{n} \sum_{x \in S} \delta(f(x) \neq h(x))$$

where  $\delta(f(x) \neq h(x))$  is 1 if  $f(x) \neq h(x)$ , and 0 otherwise

How well does  $error_S(h)$  estimate  $error_D(h)$ ?



# Problems Estimating Error

1. *Bias*: If  $S$  is training set,  $error_S(h)$  is optimistically biased

$$bias \equiv E[error_S(h)] - error_D(h)$$

For unbiased estimate,  $h$  and  $S$  must be chosen independently

2. *Variance*: Even with unbiased  $S$ ,  $error_S(h)$  may still vary from  $error_D(h)$

# Example

Hypothesis  $h$  misclassifies 12 of 40 examples in  $S$ .

$$error_S(h) = \frac{12}{40} = .30$$

What is  $error_D(h)$ ?

# Estimators

Experiment:

1. Choose sample  $S$  of size  $n$  according to distribution  $D$
2. Measure  $error_S(h)$   
 $error_S(h)$  is a random variable (i.e., result of an experiment)  
 $error_S(h)$  is an unbiased **estimator** for  $error_D(h)$

Given observed  $error_S(h)$  what can we conclude about  $error_D(h)$ ?

# Confidence Intervals

If  $S$  contains  $n$  examples,  $n \geq 30$ , drawn independently of  $h$  and each other

Then with approximately  $N\%$  probability,  $error_D(h)$  lies in interval

$$error_S(h) \pm z_N \sqrt{\frac{error_S(h)(1 - error_S(h))}{n}}$$

where

N% :	50%	68%	80%	90%	95%	98%	99%
$z_N$ :	0.67	1.00	1.28	1.64	1.96	2.33	2.53

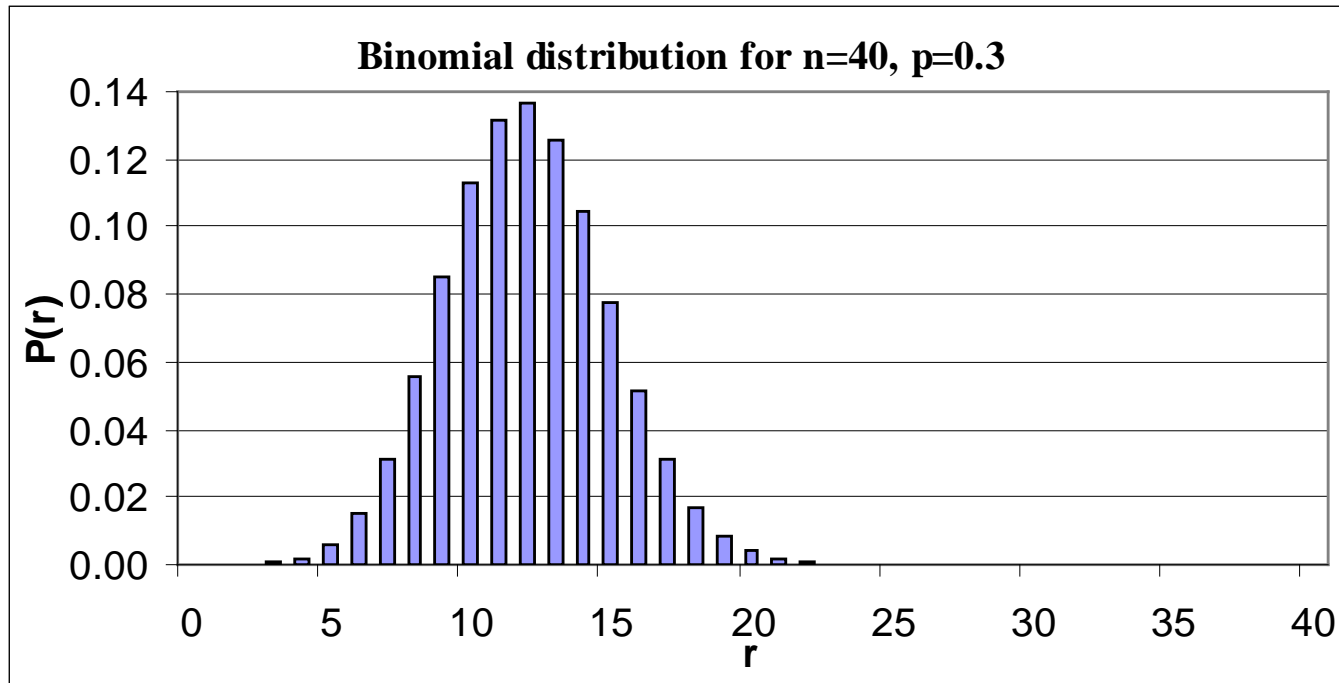
# Confidence Intervals

If  $S$  contains  $n$  examples, drawn independently of  $h$  and each other

Then with approximately 95% probability,  $error_D(h)$  lies in interval

$$error_S(h) \pm 1.96 \sqrt{\frac{error_S(h)(1 - error_S(h))}{n}}$$

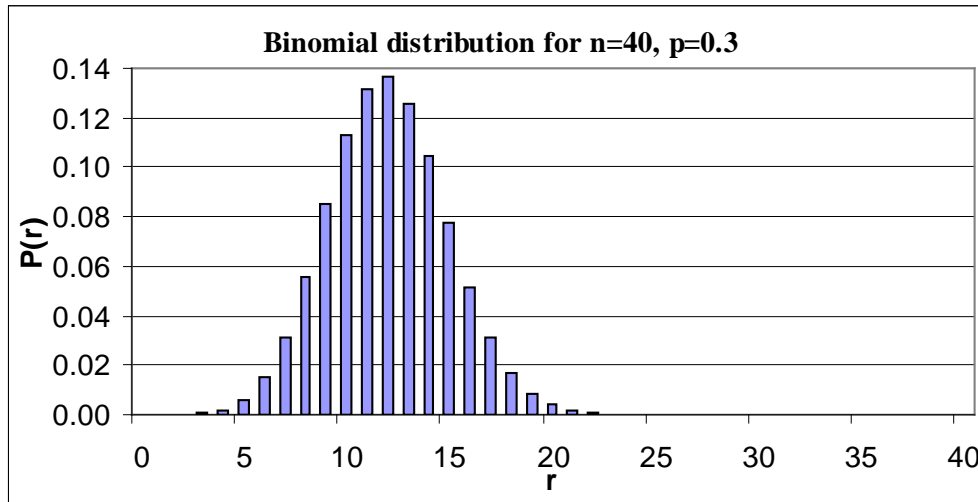
# $error_S(h)$ is a Random Variable



- ▶ Rerun experiment with different randomly drawn  $S$  (size  $n$ )
- ▶ Probability of observing  $r$  misclassified examples:

$$P(r) = \frac{n!}{r!(n-r)!} error_D(h)^r (1 - error_D(h))^{n-r}$$

# Binomial Probability Distribution

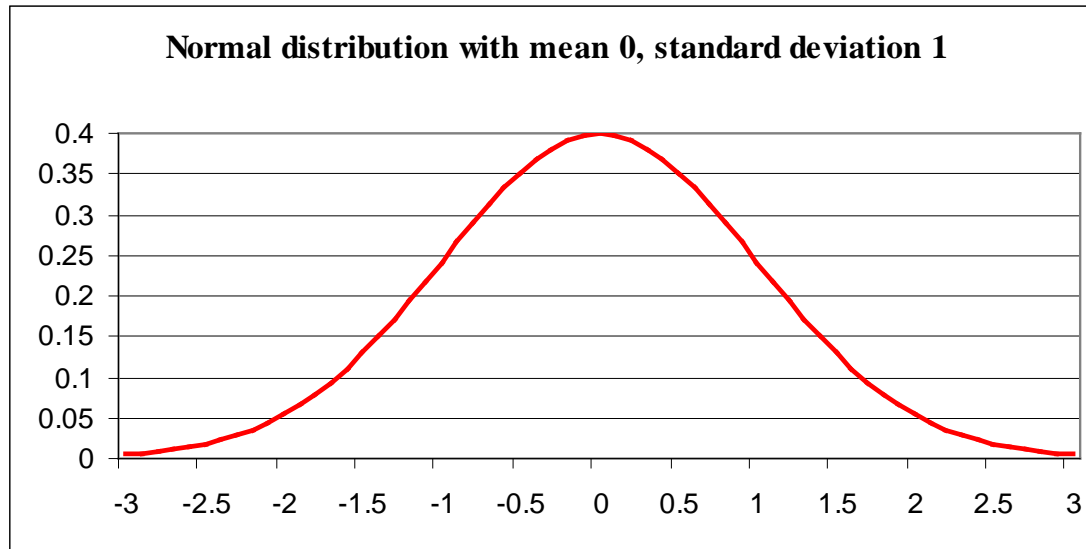


$$P(r) = \frac{n!}{r!(n-r)!} p^r (1-p)^{n-r}$$

Probability  $P(r)$  of  $r$  heads in  $n$  coin flips, if  $p = \Pr(\text{heads})$

- Expected, or mean value of  $X$  :  $E[X] \equiv \sum_{i=0}^n iP(i) = np$
- Variance of  $X$  :  $Var(X) \equiv E[(X - E[X])^2] = np(1-p)$
- Standard deviation of  $X$  :  $\sigma_X \equiv \sqrt{E[(X - E[X])^2]} = \sqrt{np(1-p)}$

# Normal Probability Distribution



$$P(r) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$$

The probability that  $X$  will fall into the interval  $(a,b)$  is given by

$$\int_a^b p(x)dx$$

- Expected, or mean value of  $X$  :  $E[X] = \mu$
- Variance of  $X$  :  $Var(X) = \sigma^2$
- Standard deviation of  $X$  :  $\sigma_X = \sigma$



# Normal Distribution Approximates Binomial

$error_s(h)$  follows a Binomial distribution, with

- mean  $\mu_{error_s(h)} = error_D(h)$
- standard deviation

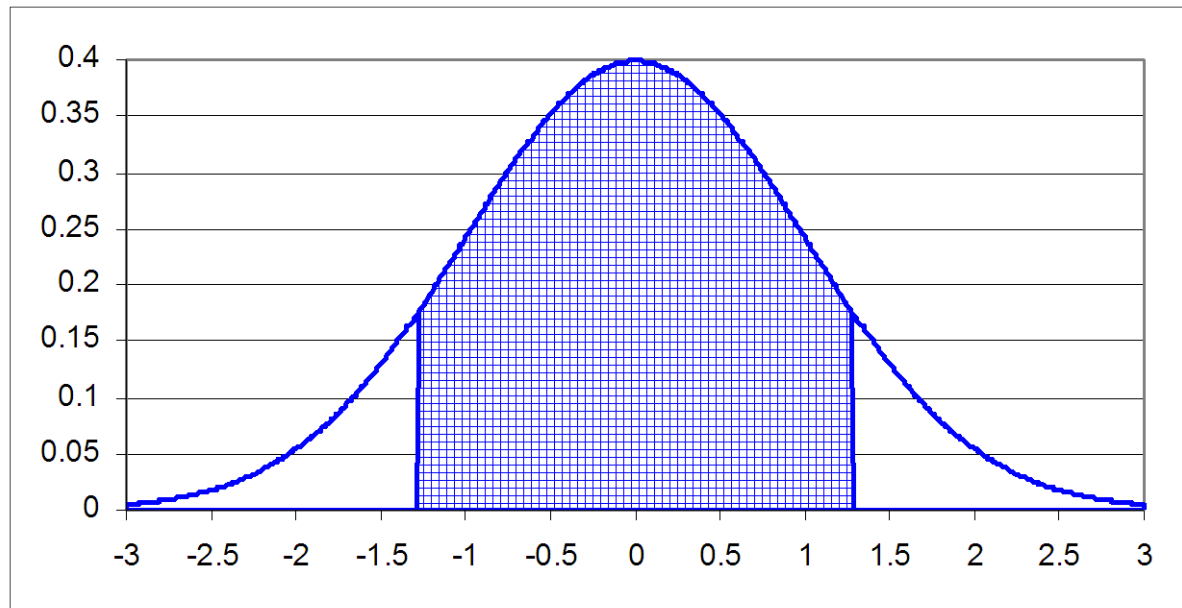
$$\sigma_{error_s(h)} = \sqrt{\frac{error_D(h)(1 - error_D(h))}{n}}$$

Approximate this by a Normal distribution with

- mean  $\mu_{error_s(h)} = error_D(h)$
- standard deviation

$$\sigma_{error_s(h)} \approx \sqrt{\frac{error_s(h)(1 - error_s(h))}{n}}$$

# Normal Probability Distribution



80% of area (probability) lies in  $\mu \pm 1.28\sigma$

N% of area (probability) lies in  $\mu \pm z_N\sigma$

N% :	50%	68%	80%	90%	95%	98%	99%
$z_N$ :	0.67	1.00	1.28	1.64	1.96	2.33	2.53

# Confidence Intervals, More Precisely

If  $S$  contains  $n$  examples, drawn independently of  $h$  and each other  $n \geq 30$

Then with approximately 95% probability,  $error_S(h)$  lies in interval

$$error_D(h) \pm 1.96 \sqrt{\frac{error_D(h)(1 - error_D(h))}{n}}$$

equivalently,  $error_D(h)$  lies in interval

$$error_S(h) \pm 1.96 \sqrt{\frac{error_D(h)(1 - error_D(h))}{n}}$$

which is approximately

$$error_S(h) \pm 1.96 \sqrt{\frac{error_S(h)(1 - error_S(h))}{n}}$$

# Calculating Confidence Intervals

1. Pick parameter  $p$  to estimate  $error_D(h)$
2. Choose an estimator  $error_S(h)$
3. Determine probability distribution that governs estimator  
 $error_S(h)$  governed by Binomial distribution,  $n \geq 30$   
approximated by Normal when
4. Find interval  $(L, U)$  such that N% of probability mass falls in the interval  
Use table of  $z_N$  values

# Central Limit Theorem

Consider a set of independent, identically distributed random variables  $Y_1 \dots Y_n$ , all governed by an arbitrary probability distribution with mean  $\mu$  and finite variance  $\sigma^2$ . Define the sample mean

$$\bar{Y} \equiv \frac{1}{n} \sum_{i=1}^n Y_i$$

**Central Limit Theorem.** As  $n \rightarrow \infty$ , the distribution governing  $\bar{Y}$  approaches a Normal distribution, with mean  $\mu$  and variance  $\frac{\sigma^2}{n}$ .

# Difference Between Hypotheses

Test  $h_1$  on sample  $S_1$ , test  $h_2$  on  $S_2$

1. Pick parameter to estimate

$$d \equiv \text{error}_D(h_1) - \text{error}_D(h_2)$$

2. Choose an estimator

$$\hat{d} \equiv \text{error}_{S_1}(h_1) - \text{error}_{S_2}(h_2)$$

3. Determine probability distribution that governs estimator

$$\sigma_d \approx \sqrt{\frac{\text{error}_{S_1}(h_1)(1 - \text{error}_{S_1}(h_1))}{n_1} + \frac{\text{error}_{S_2}(h_2)(1 - \text{error}_{S_2}(h_2))}{n_2}}$$

4. Find interval (L, U) such that N% of probability mass falls in the interval

$$\hat{d} \pm z_N \sqrt{\frac{\text{error}_{S_1}(h_1)(1 - \text{error}_{S_1}(h_1))}{n_1} + \frac{\text{error}_{S_2}(h_2)(1 - \text{error}_{S_2}(h_2))}{n_2}}$$

# Paired $t$ test to Compare $L_A, L_B$

1. Partition data into  $k$  disjoint test sets  $T_1, T_2, \dots, T_k$  of equal size, where this size is at least 30.

2. For  $i$  from 1 to  $k$  do

$$S_i = \{D - T_i\}$$

$$h_A = L_A(S_i)$$

$$h_B = L_B(S_i)$$

$$\delta_i \leftarrow \text{error}_{T_i}(h_A) - \text{error}_{T_i}(h_B)$$

3. Return the value  $d$ , where

$$\bar{\delta} \equiv \frac{1}{k} \sum_{i=1}^k \delta_i$$

$N\%$  confidence interval estimate for  $d$ :

$$\bar{\delta} \pm t_{N,k-1} s_{\bar{\delta}}$$

$$s_{\bar{\delta}} \equiv \sqrt{\frac{1}{k(k-1)} \sum_{i=1}^k (\delta_i - \bar{\delta})^2}$$

Note  $\delta_i$  approximately Normally distributed

# CV for Model selection

- ▶ **Training data set** – the set of data used to learn a model or hypothesis
- ▶ **Test data set** – the set of data used to estimate some value (often accuracy) related to a model
- ▶ **Validation set** – a set of data used to select parameters for a model, often as follows
  - Divide training data into a “sub” training set and validation set
  - For each possible set of parameters
    - Create a model using the “sub” training set
    - Evaluate the model on the validation set and pick the one that performs the best



# Bias–Variance Tradeoff: Intuition 1

- ▶ The goal in learning is not to learn an exact representation of the training data itself, but to build a statistical model of the process which generates the data. This is important if the algorithm is to have good generalization performance
- ▶ We saw that
  - models with too few parameters can perform poorly
  - models with too many parameters can perform poorly
- ▶ Need to optimize the complexity of the model to achieve the best performance
- ▶ One way to get insight into this tradeoff is the decomposition of generalization error into  $\text{bias}^2 + \text{variance}$ 
  - a model which is too simple, or too inflexible, will have a large bias
  - a model which has too much flexibility will have high variance

# Intuition

- ▶ bias:
  - measures the accuracy or quality of the algorithm
  - high bias means a poor match
- ▶ variance:
  - measures the precision or specificity of the match
  - a high variance means a weak match
- ▶ We would like to minimize each of these
- ▶ Unfortunately, we can't do this independently, there is a trade-off

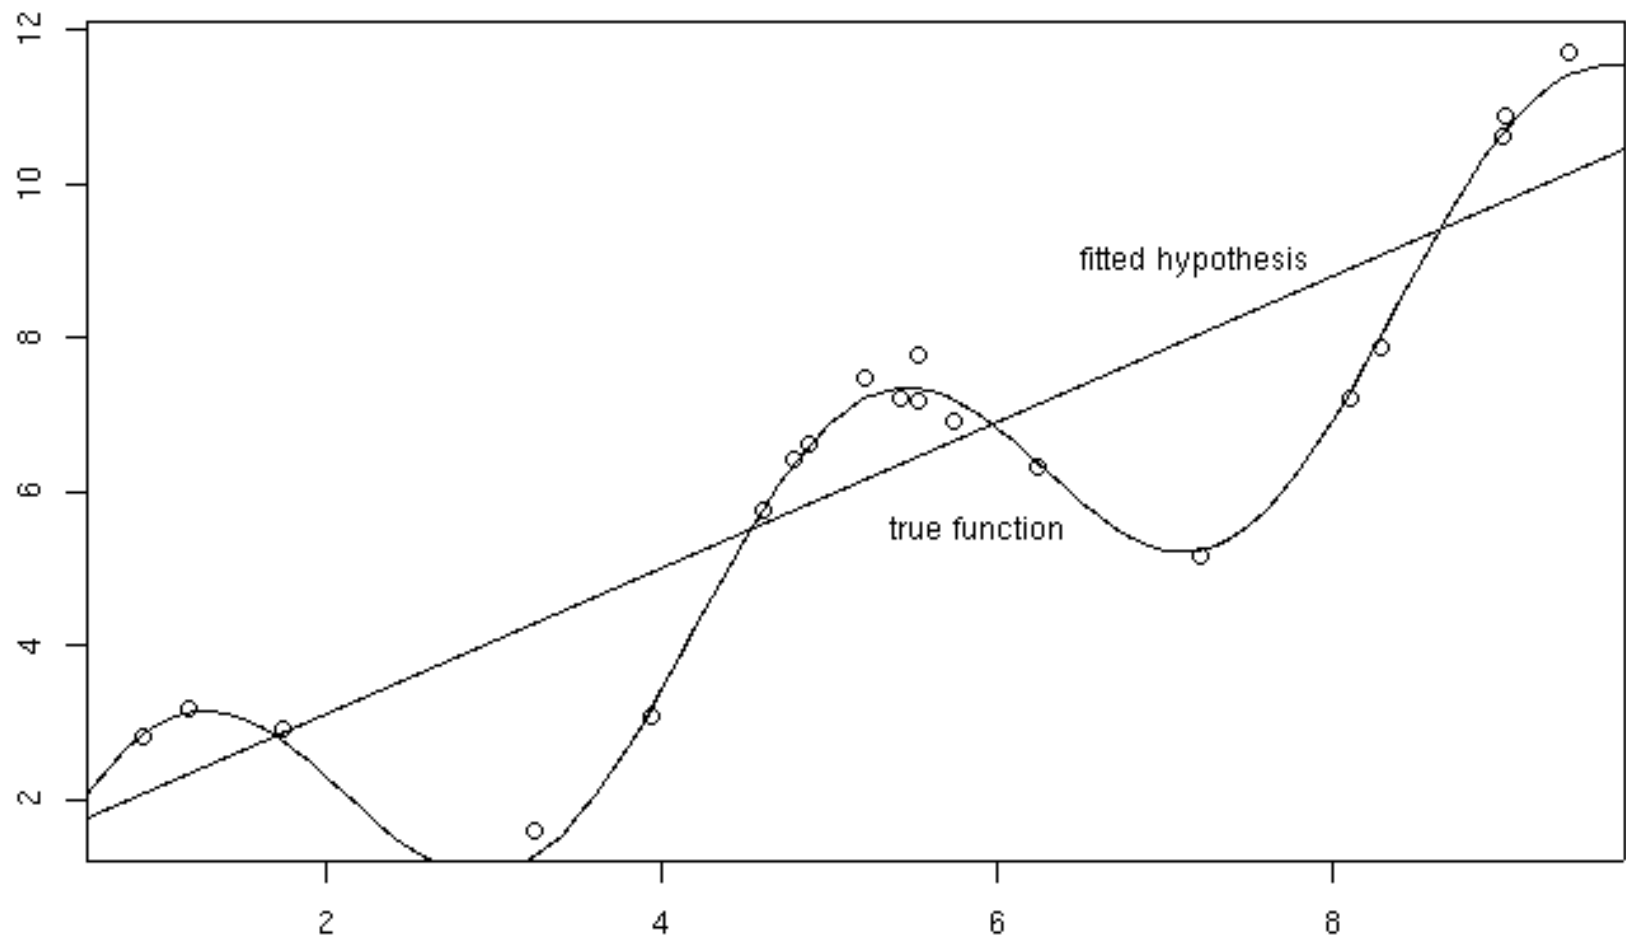
# Bias–Variance Analysis in Regression

- ▶ True function is  $y = f(x) + \varepsilon$ 
  - where  $\varepsilon$  is normally distributed with zero mean and standard deviation  $\sigma$ .
- ▶ Given a set of training examples,  $\{(x_i, y_i)\}$ , we fit an hypothesis  $h(x) = w \cdot x + b$  to the data to minimize the squared error

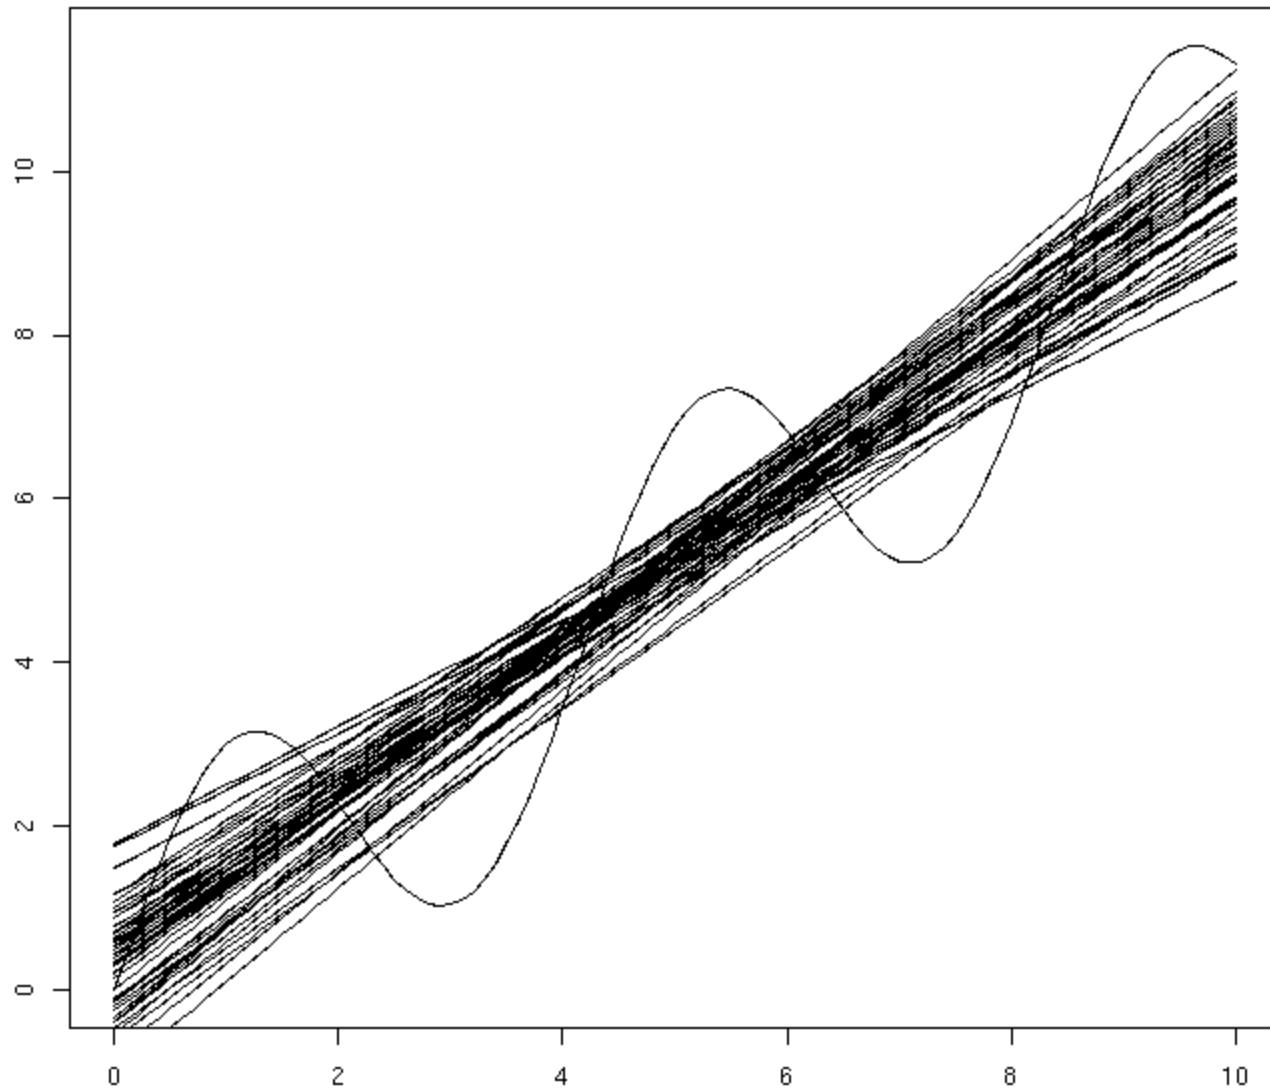
$$\sum_i [y_i - h(x_i)]^2$$

Example: 20 points

$$y = x + 2 \sin(1.5x) + N(0,0.2)$$



50 fits (20 examples each)



# Bias–Variance Analysis

- ▶ Now, given a new data point  $x^*$  (with observed value  $y^* = f(x^*) + \varepsilon$ , we would like to understand the expected prediction error

$$E[ (y^* - h(x^*))^2 ]$$

# Classical Statistical Analysis

- ▶ Imagine that our particular training sample  $S$  is drawn from some population of possible training samples according to  $P(S)$ .
- ▶ Compute  $E_p [ (y^* - h(x^*))^2 ]$
- ▶ Decompose this into “bias”, “variance”, and “noise”

# Lemma

- Let  $Z$  be a random variable with probability distribution  $P(Z)$
- Let  $\bar{Z} = E_p[ Z ]$  be the average value of  $Z$ .
- Lemma :  $E[ (Z - \bar{Z})^2 ] = E[Z^2] - \bar{Z}^2$

$$\begin{aligned} E[ (Z - \bar{Z})^2 ] &= E[ Z^2 - 2 Z \bar{Z} + \bar{Z}^2 ] \\ &= E[Z^2] - 2 E[Z] \bar{Z} + \bar{Z}^2 \\ &= E[Z^2] - 2 \bar{Z}^2 + \bar{Z}^2 \\ &= E[Z^2] - \bar{Z}^2 \end{aligned}$$

- Corollary :  $E[Z^2] = E[ (Z - \bar{Z})^2 ] + \bar{Z}^2$



# Bias–Variance–Noise Decomposition

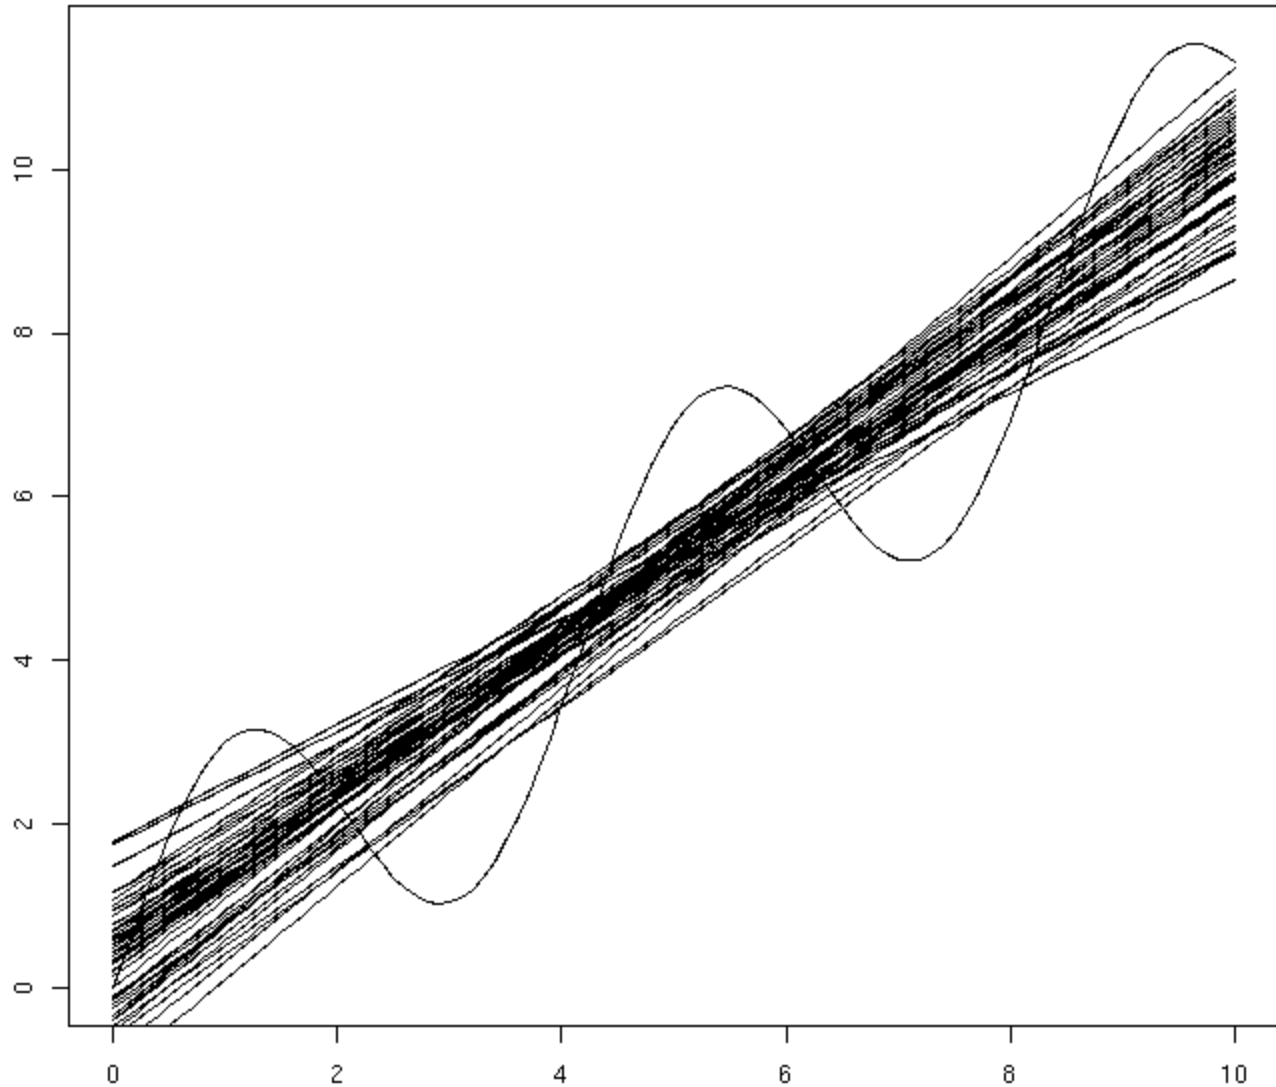
$$\begin{aligned} E[(h(x^*) - y^*)^2] &= E[h(x^*)^2 - 2h(x^*)y^* + y^{*2}] \\ &= E[h(x^*)^2] - 2E[h(x^*)]E[y^*] + E[y^{*2}] \\ &= E[(h(x^*) - \overline{h(x^*)})^2] - \overline{h(x^*)}^2 \\ &\quad - 2\overline{h(x^*)}f(x^*) \\ &\quad + E[(y^* - f(x^*))^2] + f(x^*)^2 \\ &= E[(h(x^*) - \overline{h(x^*)})^2] + \text{VARIANCE} \\ &\quad \left( \overline{h(x^*)}^2 - f(x^*) \right)^2 \text{BIAS} \\ &\quad + E[(y^* - f(x^*))^2] \text{NOISE} \\ &= \text{Var}(h(x^*)) + \text{Bias}(h(x^*))^2 + E[\epsilon^2] \\ &= \text{Var}(h(x^*)) + \text{Bias}(h(x^*))^2 + \sigma^2 \end{aligned}$$

Expected prediction error = Variance + Bias<sup>2</sup> + Noise<sup>2</sup>

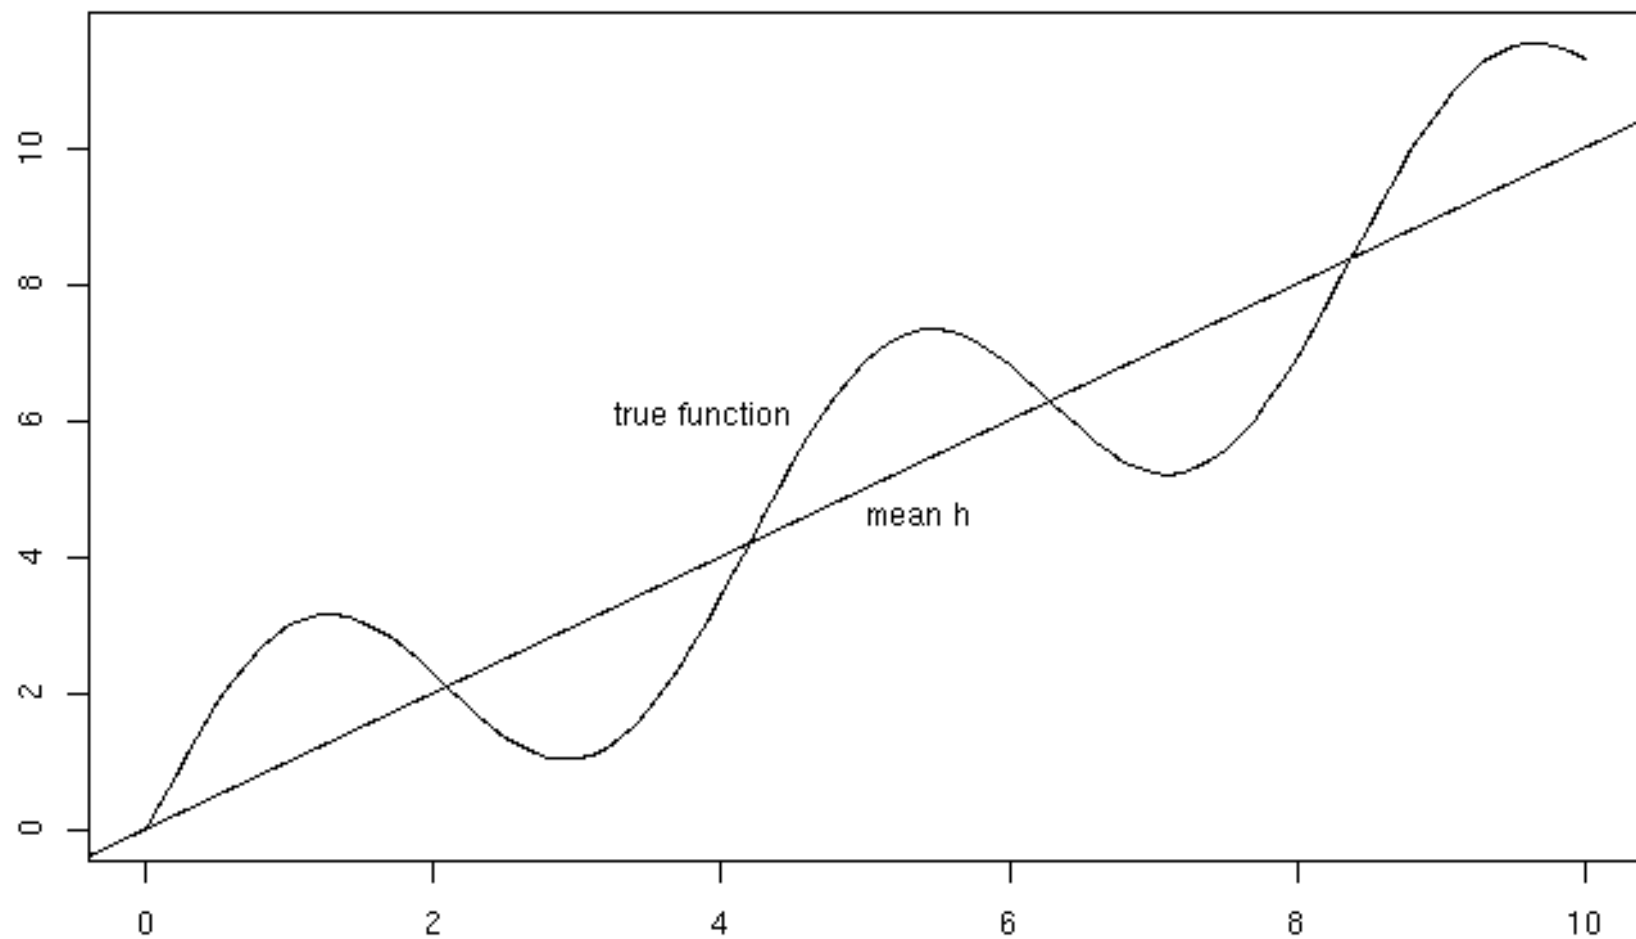
# Bias, Variance, and Noise

- ▶ Variance:  $E[ (h(x^*) - \overline{h(x^*)})^2 ]$   
Describes how much  $h(x^*)$  varies from one training set  $S$  to another
- ▶ Bias:  $[\overline{h(x^*)} - f(x^*)]$   
Describes the average error of  $h(x^*)$ .
- ▶ Noise:  $E[ (y^* - f(x^*))^2 ] = E[\varepsilon^2] = \sigma^2$   
Describes how much  $y^*$  varies from  $f(x^*)$

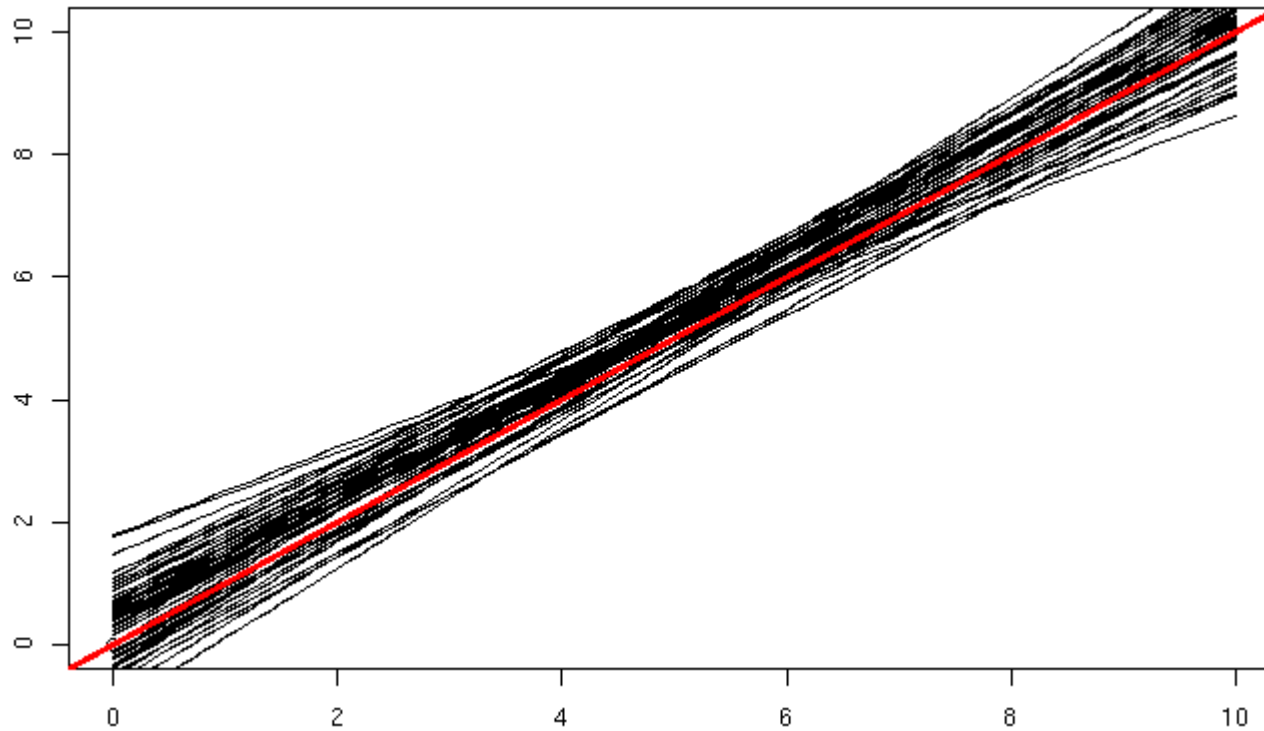
# 50 fits (20 examples each)



# Bias



# Variance



# Noise

