

CMSC726

Lecture 14

Computational Learning Theory

Ben London

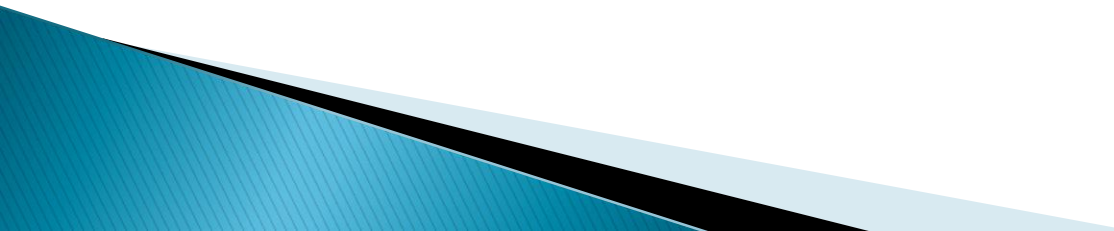
October 21, 2010




# Introduction

- ▶ What is *computational learning theory* (CLT)?
  - Complexity and correctness of machine learning
    - Time complexity
    - Sample complexity
    - Generalization error
    - Robustness (to noise)
    - Hardness

# Introduction

- ▶ Why should I care about CLT?
    - Having an understanding of learning theory will help you better understand the machine learning
    - Theory can lead to/explain/give bounds for application
    - Theory is fun! (?)
- 

# Introduction

- ▶ Goals for this lecture:
    - Review basic learning theory concepts
    - Discuss two classic learning paradigms
    - Quantify hypothesis complexity
    - Bound generalization error
    - Inspire you to get interested in CLT!
- 

# Terminology

## ▶ Instance space:

- Domain (input) of problem
- Notation:  $X$
- Can be boolean, integral, real-valued or categorical
  - CLT usually considers boolean
    - Simpler
    - Can represent any other type
- e.g. boolean strings of length- $n$

$$X \in \{0, 1\}^n$$

- size of instance space  $|X| = 2^n$

## ▶ Instance:

- a.k.a **example, sample**
- A  $n$ -tuple permutation of the instance space

# Terminology


## ► Concept:

- A deterministic partitioning (labeling) of the instance space
- A function of the instance space
- Is unique
- Notation:  $c(x)$

## ► Concept class:

- An abstraction of related concepts
- e.g. monotone DNFs, CNFs, half-spaces in  $\mathbb{R}^d$
- Notation:  $\mathcal{C}$

- Note difference: similar to object vs. class in OOP



# Terminology

## ▶ Hypothesis:

- The learning algorithm's approximation (estimate) to the target concept  $c$
- Notation: by  $h(x)$

## ▶ Hypothesis space:

- Space of all possible hypotheses for concept class  $C$
- Notation:  $H$

## ▶ Similar to concept/concept class relationship

# Terminology

## ▶ Example Oracle:


- Like a teacher, provides examples of a given concept
- Notation:
  - Malicious:  $Ex(c)$
  - Agnostic:  $Ex(c,D)$



# General Learning Paradigm

- ▶ Given:
  - Instance space  $X$
  - Concept class  $C$  over  $X$
  - Learning algorithm (learner)  $A$  for  $C$
  - (supervised) Oracle for concept  $c$  in  $C$ ,  $Ex(c)$  (or  $Ex(c, D)$ )
- ▶ After  $m$  examples, learner has hypothesis  $h$  in  $H$ , such that  $h$  is either consistent with  $c$ , or within tolerance

# Online Learning

- ▶ Often referred to as the *Online Mistake Bound* (OLMB) Model
  - ▶ Sequential, continual learning
  - ▶ Learning proceeds as a sequence of trials
  - ▶ “Malicious” Oracle  $Ex(c)$ : provides sequence of labeled examples  $(x, c(x))$  in presumably worst possible order for learning concept  $c$
- 

# Online Learning

## ▶ Given:

- Instance space  $X$ , concept class  $C$
- Learning algorithm  $A$  for concept class  $C$
- “Malicious” Oracle  $Ex(c)$  for target concept  $c$

## ▶ Trial:

1. Learner gets labeled example  $(x, c(x))$  from Oracle
2. Learner outputs  $h(x)$  using current hypothesis
3. If  $h(x) \neq c(x)$ , learner incurs a mistake
4. Learner updates hypothesis given outcome

# Online Learning

## ▶ Mistake Bound:

- Algorithm  $A$  has mistake bound  $M$  if, for any target concept  $c$  in  $C$  and any sequence of examples from  $Ex(c)$ ,  $A$  makes at most  $M$  mistakes
- After  $M$  mistakes, the learner will have hypothesis  $h$  that is consistent with  $c$
- In other words, mistake bound is a measure of algorithm convergence

# Online Learning

## ▶ Theorem:

- For any finite hypothesis space  $H$  for concept class  $C$ , there exists an algorithm with mistake bound at most  $\log|H|$

## ▶ Halving Algorithm:

1. Initialize “working” hypothesis space  $H'$  to  $H$
2. At each trial:
  1. Learner gets labeled example  $(x, c(x))$  from Oracle
  2. Learner makes prediction  $h(x)$  based on majority vote of all hypotheses in  $H'$
  3. If  $h(x) = c(x)$ , continue
  4. Else, eliminate all hypotheses from  $H'$  that predicted consistent with  $h(x)$

# Online Learning

- ▶ Halving algorithm has good mistake bound but terrible performance

- e.g.  $H$  is all monotone disjunctions of length- $n$

$$x_1 \vee x_2 \vee \cdots \vee x_n$$
$$|H| = 2^n$$

- Mistake bound  $M \leq n$
  - However, time complexity of the algorithm is  $O(2^n)$

- ▶ Fact:

- If a concept class is efficiently learnable in the online mistake bound model, it is efficiently learnable in other models

# PAC Learning

- ▶ *Probably Approximately Correct* (PAC) Model
- ▶ Batch learning: learner trains on a random set of i.i.d. examples  $S$  drawn from distribution  $D$
- ▶ Agnostic Oracle  $E(c, D)$ : provides random draws of labeled examples  $(x, c(x))$  from distribution  $D$
- ▶ Learning is approximate: we allow hypothesis to have error  $\epsilon$ , for  $0 \leq \epsilon \leq 0.5$
- ▶ We allow learner to fail with probability  $\delta$ , for  $0 \leq \delta \leq 0.5$

# PAC Learning

## ▶ Given:

- Distribution  $D$  over instance space  $X$
- Learning algorithm  $A$  for concept class  $C$
- Parameters  $\epsilon$  and  $\delta$
- Agnostic (Randomized) Oracle  $Ex(c, D)$

## ▶ Goal:

- if  $A$  is given  $m$  examples from  $Ex(c, D)$ , then, with probability  $\geq 1 - \delta$ ,  $A$  outputs hypothesis  $h$  with error  $P_D[h(x) \neq c(x)] \leq \epsilon$



# PAC Learning

- ▶ Sample complexity:
  - Number of draws from  $Ex(c, D)$  necessary to learn  $h$  within bounds of given parameters,  $\epsilon$  and  $\delta$
  - Notation:  $m$
  - We want a bound on  $m$ , given  $\epsilon$  and  $\delta$
  - Corollary: given  $m$  and  $\delta$ , we can bound  $\epsilon$

# PAC Learning Bounds

## Finite Hypothesis Spaces

- ▶ Consistent Hypothesis Finder (CHF):
  - For any sequence  $S$  of  $m$  examples labeled according to  $c$  in  $C$ , finds a consistent hypothesis  $h$  from finite hypothesis space  $H$
- ▶ Theorem:
  - Given a CHF for  $C$ , can PAC learn any  $c$  in  $C$  with sample complexity

$$m = \frac{1}{\epsilon} \left( \ln |H| + \ln \frac{1}{\delta} \right)$$

# PAC Learning Bounds

## Finite Hypothesis Spaces

### ► Corollary:

- If any hypothesis  $h$ , from a finite hypothesis space  $H$ , is consistent with  $m$  examples drawn from distribution  $D$ , its error w.r.t.  $D$  can be bound by

$$\Pr_D[h(x) \neq c(x)] \leq \frac{1}{m} \left( \ln |H| + \ln \frac{1}{\delta} \right)$$

# PAC Learning Bounds

## Finite Hypothesis Spaces

### ▶ Imperfect Hypotheses

- Suppose  $h$  is not perfect w.r.t.  $S$ , but has training error  $\epsilon_T$
- We have the following error bound

$$\Pr_D[h(x) \neq c(x)] \leq \epsilon_T + \sqrt{\frac{\ln |H| + \ln \frac{1}{\delta}}{2m}}$$

- Follows intuition: true error will be greater than the training error

# Infinite Hypothesis Spaces

- ▶ But what about infinite hypothesis spaces?
  - e.g. half-spaces in  $\mathbb{R}^d$
- ▶ How do we quantify  $|H|$ ?
- ▶ If we can't count the hypotheses directly, perhaps we can count the number of examples that can be discriminated

# Shattering





- ▶ Definition: Given a set  $S$  of  $n$  points from  $X$ , if for any labeling according to  $C$  there exists some  $h$  in  $H$  that induces this labeling, then  $H$  shatters  $S$
  
- ▶ Procedure:
  - Given:  $X, C, H, S, n$
  - 1. Consider all possible configuration of labels over  $S$   
(for simplicity, just we'll use  $+/-$  labels)
  - 2. For each configuration, verify that there exists a classifier  $h$  in  $H$  that induces it

# Shattering

## ▶ Example 1:


- $X = 1D$  number line
- $C = +/-$  labels

- $H = \text{point (separator)}$
- $n = 2$



# Shattering

- ▶ How about 3 points?



Not  
shatterable!




# Shattering


## ▶ Example 2:

- $X = 2D$  plane
- $C = +/-$  labels


- $H = \text{line}$
- $n = 3$



Isolate any 1 point





Isolate any 2 points



Isolate any 3 points

# Shattering

- ▶ How about 4 points?



Not  
shatterable!

# VC Dimension

- ▶ Vapnik–Chervonenkis Dimension
- ▶ Determines capacity of a (possibly infinite) hypothesis space
- ▶ Definition: the largest number of points that can be shattered by hypothesis space  $H$
- ▶ Notation:  $\text{VCDim}(H)$ 
  - Alt  $\text{VCDim}(C)$
- ▶ Examples:
  - $C = +/ -$  labels of points in 1-D  $\text{VCDim}(H) = 2$
  - $C = +/ -$  labels of points in 2-D  $\text{VCDim}(H) = 3$
  - $C = +/ -$  labels of points in  $n$ -D  $\text{VCDim}(H) = n + 1$

# VC Dimension

- ▶ Typical proof:
  - Given: concept class  $C$ , hypothesis space  $H$ 
    1. Lower bound: show that there exists some set of at least  $n$  points that can be shattered by  $H$
    2. Upper bound: show that there does not exist any set of  $n+1$  points that can be shattered by  $H$
  - Proof is usually analytic and/or geometric

# VC Dimension

- ▶ VC Dimension gives a “reasonable” estimate of the size (capacity, expressivity) of infinite hypothesis spaces
- ▶ If we can induce  $2^{\text{VCDim}(H)}$  labelings, then there are at least  $2^{\text{VCDim}(H)}$  possible hypotheses

# VC Dimension

## Simple PAC Bounds for Binary Labels

► Note:  $\ln 2^{\text{VCDim}(H)} = O(\text{VCDim}(H))$

► Sample complexity:

$$m = \frac{1}{\epsilon} \left( O(\text{VCDim}(H)) + \ln \frac{1}{\delta} \right)$$

► Error bound:

$$\Pr_D[h(x) \neq c(x)] \leq \epsilon_T + \sqrt{\frac{O(\text{VCDim}(H)) + \log \frac{1}{\delta}}{2m}}$$

# VC Dimension

## General PAC Bounds

- ▶ Sample complexity:

$$m = O\left(\frac{d}{\epsilon} \log \frac{1}{\epsilon} + \frac{1}{\epsilon} \log \frac{1}{\delta}\right) \quad d = VCDim(H)$$

- ▶ Error bound:

$$\Pr_D[h(x) \neq c(x)] \leq \epsilon_T + \sqrt{\frac{d \left(\log \frac{2m}{d} + 1\right) + \log \frac{4}{\delta}}{m}}$$

# VC Dimension

## ▶ Warnings:

- VC dimension can be *infinite*
  - If VC dimension is infinite, then no finite set of examples is enough to PAC learn a concept class
- VC dimension is *distribution free*: it makes no assumptions about the underlying distribution of the data
  - Bounds can be overly pessimistic (but still useful)



# Interpretation

- ▶ Expressivity vs. Generality:
  - A more expressive model can capture more detail, but it may also overfit the training data
  - A more general model is less likely to overfit, but it can't capture as much detail
- ▶ From VC Dimension perspective:
  - Increasing  $VCDim(H)$  may decrease empirical error but increase generalization error
  - Decreasing  $VCDim(H)$  may decrease generalization error but increase empirical error
- ▶ Empirical risk vs. true risk minimization

# Risk Minimization

- ▶ Risk = Error

- Expectation of a loss function  $L(y, x, \theta)$

$$R(\theta) = \mathbb{E}[L(y_i, x_i, \theta)] = \int_X \int_Y \Pr[x, y] L(y, x, \theta) dx dy$$

- ▶ But we don't know  $P(X, Y)$ !

- (Unless we have infinite data)

- ▶ Best we can do is approximate true risk

⇒ Empirical risk

# Empirical Risk Minimization


- ▶ Most learning algorithms minimize empirical risk:

$$\hat{R}(\theta) = \frac{1}{m} \sum_{i=1}^m L(y_i, x_i, \theta)$$

- Quadratic:  $L(y, x, \theta) = \frac{1}{2} (y - f(x; \theta))^2$  Regression
- Binary:  $L(y, x, \theta) = \text{sign}(-yf(x; \theta))$  Perceptron

# Empirical Risk Minimization

- ▶ Problem: empirical risk is not true risk



# Structural Risk Minimization

- ▶ We want a guaranteed risk  $J(\hat{\theta})$
- ▶ Solution: add confidence bound to empirical risk  $J(\theta) = \hat{R}(\theta) + C(\theta)$
- ▶ If  $\forall \theta, R(\theta) \leq J(\theta)$ , then we have true risk bound
- ▶ If we select the  $\hat{\theta} = \operatorname{argmin}_{\theta} J(\theta)$ , then we can guarantee that the true error  $R(\hat{\theta}) \leq J(\hat{\theta})$

# Structural Risk Minimization

- ▶ Recall the PAC error bound

$$\Pr_D[h(x) \neq c(x)] \leq \epsilon_T + \sqrt{\frac{d \left( \log \frac{2m}{d} + 1 \right) + \log \frac{4}{\delta}}{m}}$$

- ▶ This gives us

True risk

$$R(\theta) = \Pr_D[h(x) \neq c(x)]$$

Empirical risk


$$\hat{R}(\theta) = \epsilon_T$$

Model complexity

$$C(\theta) = \sqrt{\frac{d \left( \log \frac{2m}{d} + 1 \right) + \log \frac{4}{\delta}}{m}}$$

# Structural Risk Minimization

- ▶ Bounding true risk



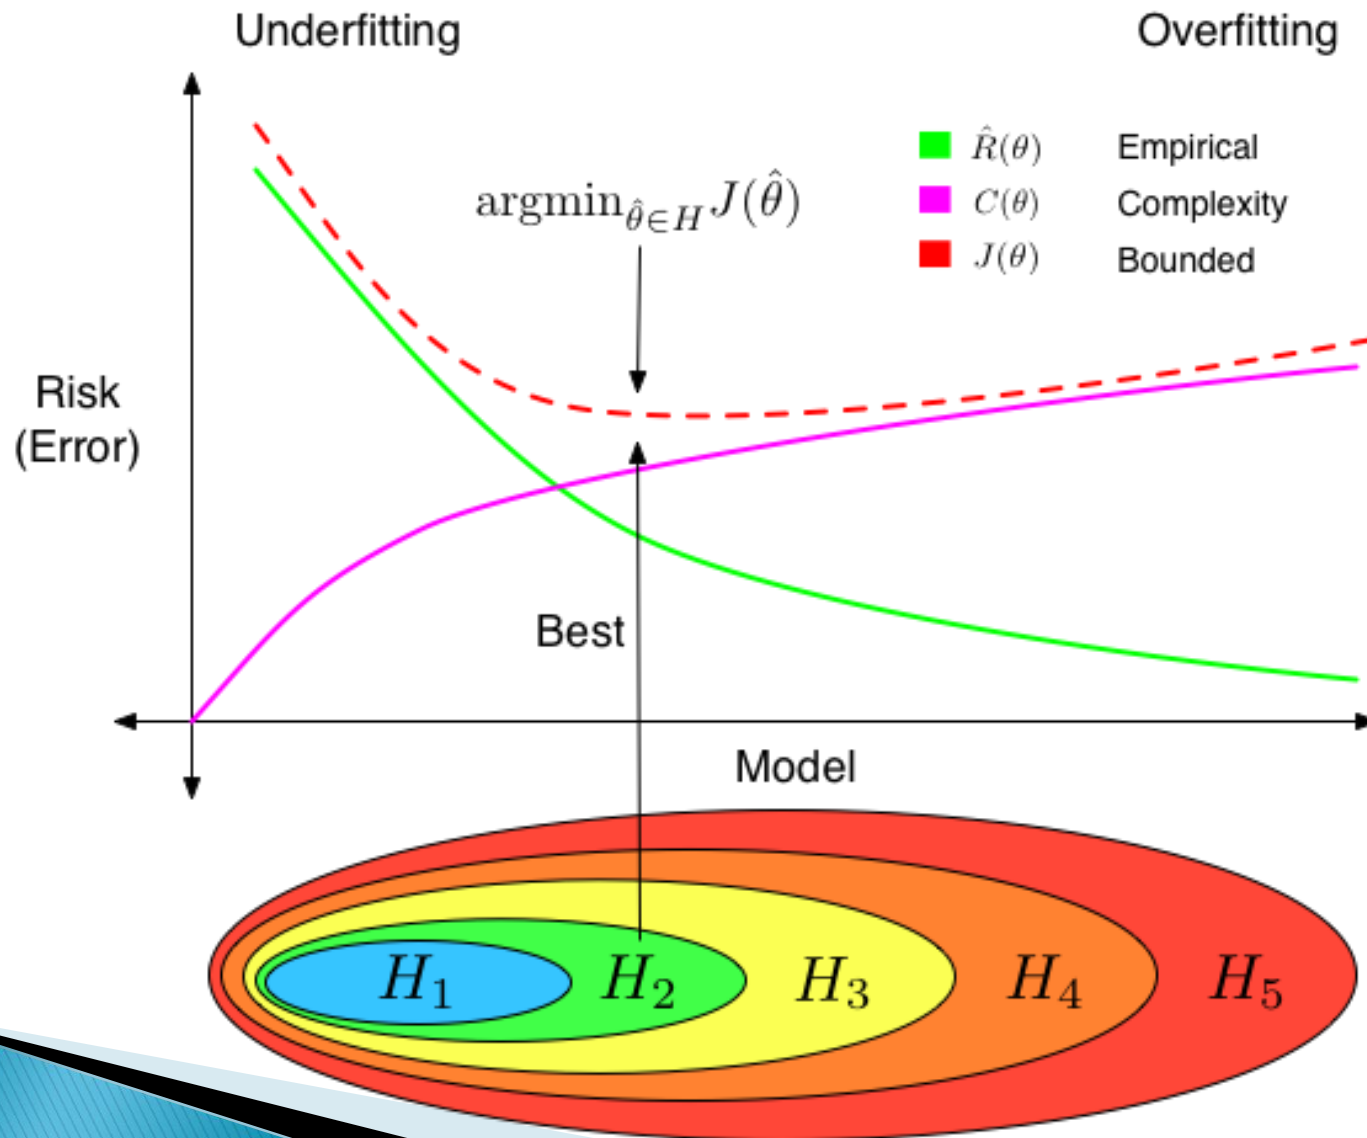
# Structural Risk Minimization

- ▶ Goal: choose the hypothesis that minimizes error bound  $J(\theta)$
- ▶ Procedure:
  1. Select set of “reasonable” hypothesis spaces using *a priori* knowledge of the problem
  2. For each hypothesis space:
    1. Train hypothesis that minimizes empirical error
    2. Compute error bound  $J(\hat{\theta}) = \hat{R}(\hat{\theta}) + C(\hat{\theta})$
  3. Select hypothesis that minimizes error bound

$$\theta^* = \operatorname{argmin}_{\hat{\theta}} J(\hat{\theta})$$



# Structural Risk Minimization




# Philosophical Interpretation: Occam's Razor

- ▶ “Entities must not be multiplied beyond necessity”
  - William of Ockham (14<sup>th</sup> cent. AD)
- ▶ Given many potential hypotheses, choose the one that makes the fewest assumptions




# VC Dimension & Large Margins

- ▶ Arbitrary linear classifiers are too flexible
- ▶ Can reduce VC dimension if we restrict them
- ▶ Reasonable assumption: constrain data to living inside a circle/sphere
  - dimension  $d$
  - diameter  $D$
- ▶ Apply linear classifier with margin  $M$




# VC Dimension & Large Margins

- ▶ If  $M$  is small, can shatter 3 points




# VC Dimension & Large Margins

- ▶ If  $M$  is large enough, can only shatter 2 points



Can't shatter

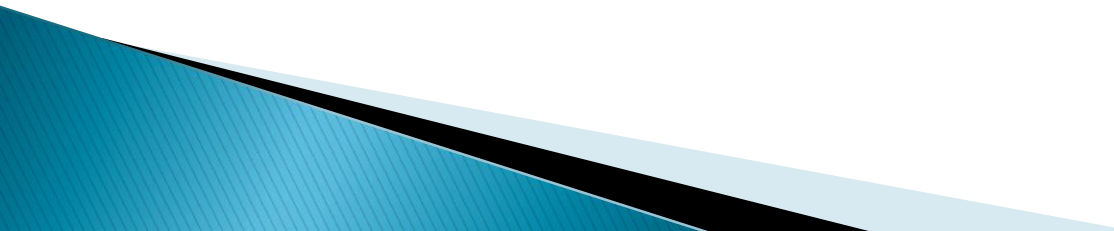


Can shatter

# VC Dimension & Large Margins

- ▶ For hyperplanes, as  $M$  grows relative to  $D$ , VC dimension goes down
- ▶ General formula  $\text{VCDim}(H) = \min \left( \left\lceil \frac{D^2}{M^2} \right\rceil, d \right) + 1$ 
  - Arbitrary data:  $\text{VCDim}(H) \rightarrow d + 1$
  - Inside sphere:  $\text{VCDim}(H) \rightarrow \lceil D^2/M^2 \rceil + 1$ 
    - Typical data lives inside sphere
- ▶  $\therefore$  Larger margins decrease VC dimension  
 $\Rightarrow$  Decrease generalization error!

# Relation to SVMs

- ▶ SVMs maximize margin,  
⇒ minimize generalization error
  - ▶ SVMs naturally perform SRM
  - ▶ CLT explains why SVMs are so good
- 

# Research Topics in CLT

- ▶ Learning models
    - Online
    - PAC
    - Active
    - Statistical Query
    - Boolean Hypercube
    - Evolvability
  - ▶ Cryptography/hardness results
  - ▶ Boosting/ensemble methods
- 